

Porter- Neural Networks Regression

- Porter, India's largest marketplace for intra-city logistics, is revolutionizing the delivery sector with technology-driven solutions.
- Task: Given new order from customer, Help Porter to estimate delivery time, for that we build NN regression model.

- Each row in this file corresponds to one unique delivery. Each column corresponds to a feature as explained below.

| Sr. No. | Feature Name | Description |
|---------|--------------------------|--|
| 1 | market_id | An integer ID indicating the market area of the restaurant. |
| 2 | created_at | Timestamp of when the order was placed. |
| 3 | actual_delivery_time | Timestamp of when the order was delivered. |
| 4 | store_id | Store ID |
| 5 | store_primary_category | Category classification of the restaurant. |
| 6 | order_protocol | Numeric code representing the mode of order placement. (how the order was placed ie: through porter, call to restaurant, pre booked, third part etc) |
| 7 | total_items | Total items in order |
| 8 | subtotal | A combined feature detailing the total number of items and final price of the order before taxes and fees. |
| 9 | num_distinct_items | Count of different items in the order. |
| 10 | min_item_price | Price of the least expensive item in the order. |
| 11 | max_item_price | Price of the most expensive item in the order. |
| 12 | total_onshift_partners | Number of delivery partners on duty when the order was placed. |
| 13 | total_busy_partners | Number of delivery partners busy with other tasks at the order placement time. |
| 14 | total_outstanding_orders | Total count of orders pending at the time. |

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import warnings
4 warnings.filterwarnings("ignore")
5 import seaborn as sns
6 import matplotlib.pyplot as plt
7
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.model_selection import train_test_split
10 from category_encoders import TargetEncoder
11 from sklearn.metrics import mean_squared_error, r2_score
12
13 import tensorflow as tf
14 import keras
15 from tensorflow.keras.models import Sequential
16 from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
17 from tensorflow.keras.callbacks import EarlyStopping
18 from tensorflow.keras.optimizers import Adam
19 import keras_tuner
```

In [2]:

```
1 # data Loading as df
2 df= pd.read_csv("Porter.csv")
```

In [3]:

1df.head()

Out[3]:

| | market_id | created_at | actual_delivery_time | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_partners |
|---|-----------|---------------------|----------------------|----------------------------------|------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|---------------------|
| 0 | 1.0 | 2015-02-06 22:24:17 | 2015-02-06 23:27:16 | df263d996281d984952c07998dc54358 | american | 1.0 | 4 | 3441 | 4 | 557 | 1239 | 33.0 | |
| 1 | 2.0 | 2015-02-10 21:49:25 | 2015-02-10 22:56:29 | f0ade77b43923b38237db569b016ba25 | mexican | 2.0 | 1 | 1900 | 1 | 1400 | 1400 | 1.0 | |
| 2 | 3.0 | 2015-01-22 20:39:28 | 2015-01-22 21:09:09 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | 1 | 1900 | 1 | 1900 | 1900 | 1.0 | |
| 3 | 3.0 | 2015-02-03 21:21:45 | 2015-02-03 22:13:00 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | 6 | 6900 | 5 | 600 | 1800 | 1.0 | |
| 4 | 3.0 | 2015-02-15 02:40:36 | 2015-02-15 03:20:26 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | 3 | 3900 | 3 | 1100 | 1600 | 6.0 | |

- Data has "actual_delivery_time" and "created_at" which is important to calculate total time taken to deliver.
- Data also have "total_items" and "num_distinct_items" which will be helpful to determine if more number of item in single order taking significantly more time or not.

In [4]:

1df.shape

Out[4]: (197428, 14)

In [5]:

1df.columns

Out[5]: Index(['market_id', 'created_at', 'actual_delivery_time', 'store_id', 'store_primary_category', 'order_protocol', 'total_items', 'subtotal', 'num_distinct_items', 'min_item_price', 'max_item_price', 'total_onshift_partners', 'total_busy_partners', 'total_outstanding_orders'], dtype='object')

In [6]:

1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 197428 entries, 0 to 197427
Data columns (total 14 columns):
Column Non-Null Count Dtype
--- -
0 market_id 196441 non-null float64
1 created_at 197428 non-null object
2 actual_delivery_time 197421 non-null object
3 store_id 197428 non-null object
4 store_primary_category 192668 non-null object
5 order_protocol 196433 non-null float64
6 total_items 197428 non-null int64
7 subtotal 197428 non-null int64
8 num_distinct_items 197428 non-null int64
9 min_item_price 197428 non-null int64
10 max_item_price 197428 non-null int64
11 total_onshift_partners 181166 non-null float64
12 total_busy_partners 181166 non-null float64
13 total_outstanding_orders 181166 non-null float64
dtypes: float64(5), int64(5), object(4)
memory usage: 21.1+ MB

- created_at and actual_delivery_time given as object but it should be as timestamp.
- Some feature consist of numerical values and some are string

In [7]:

1 df["created_at"] = pd.to_datetime(df["created_at"])
2 df["actual_delivery_time"] = pd.to_datetime(df["actual_delivery_time"])

In [8]:

1 df.describe()

Out[8]:

| | market_id | created_at | actual_delivery_time | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_partners | total_outstanding_orders |
|-------|---------------|-------------------------------|-------------------------------|----------------|---------------|---------------|--------------------|----------------|----------------|------------------------|---------------------|--------------------------|
| count | 196441.000000 | 197428 | 197421 | 196433.000000 | 197428.000000 | 197428.000000 | 197428.000000 | 197428.000000 | 197428.000000 | 181166.000000 | 181166.000000 | 181166.000000 |
| mean | 2.978706 | 2015-02-04 22:00:09.537962752 | 2015-02-04 22:48:23.348914432 | 2.882352 | 3.196391 | 2682.331402 | 2.670791 | 686.218470 | 1159.588630 | 44.808093 | 41.739747 | 58.050065 |
| min | 1.000000 | 2014-10-19 05:24:15 | 2015-01-21 15:58:11 | 1.000000 | 1.000000 | 0.000000 | 1.000000 | -86.000000 | 0.000000 | -4.000000 | -5.000000 | -6.000000 |
| 25% | 2.000000 | 2015-01-29 02:32:42 | 2015-01-29 03:22:29 | 1.000000 | 2.000000 | 1400.000000 | 1.000000 | 299.000000 | 800.000000 | 17.000000 | 15.000000 | 17.000000 |
| 50% | 3.000000 | 2015-02-05 03:29:09.500000 | 2015-02-05 04:40:41 | 3.000000 | 3.000000 | 2200.000000 | 2.000000 | 595.000000 | 1095.000000 | 37.000000 | 34.000000 | 41.000000 |
| 75% | 4.000000 | 2015-02-12 01:39:18.500000 | 2015-02-12 02:25:26 | 4.000000 | 4.000000 | 3395.000000 | 3.000000 | 949.000000 | 1395.000000 | 65.000000 | 62.000000 | 85.000000 |
| max | 6.000000 | 2015-02-18 06:00:44 | 2015-02-19 22:45:31 | 7.000000 | 411.000000 | 27100.000000 | 20.000000 | 14700.000000 | 14700.000000 | 171.000000 | 154.000000 | 285.000000 |
| std | 1.524867 | NaN | NaN | 1.503771 | 2.666546 | 1823.093688 | 1.630255 | 522.038648 | 558.411377 | 34.526783 | 32.145733 | 52.661830 |

- we have order created data from 2015-02-04 to 2015-02-18.
- Maximum total item in single delivery is 411
- Maximum subtotal is 27100 but 75% have subtotal less than equal to 3395
- Maximum Number of distinct Item in single delivery is 20.
- 50% of the time there are less than or equal to 2 items in each delivery

```
In [9]: 1 df.describe(include= "object")
```

Out[9]:

| | store_id | store_primary_category |
|--------|----------------------------------|------------------------|
| count | 197428 | 192668 |
| unique | 6743 | 74 |
| top | d43ab110ab2489d6b9b2caa394bf920f | american |
| freq | 937 | 19399 |

- Most frequent store is have store_id "d43ab110ab2489d6b9b2caa394bf920f"
- Most frequent category of restaurant is "american"

```
In [10]: 1 df.isna().sum()
```

```
Out[10]: market_id          987
created_at              0
actual_delivery_time    7
store_id                0
store_primary_category  4760
order_protocol          995
total_items             0
subtotal               0
num_distinct_items      0
min_item_price          0
max_item_price          0
total_onshift_partners  16262
total_busy_partners     16262
total_outstanding_orders 16262
dtype: int64
```

```
In [11]: 1 # In percentage
2 df.isna().sum()/ len(df)*100
```

```
Out[11]: market_id          0.499929
created_at              0.000000
actual_delivery_time    0.003546
store_id                0.000000
store_primary_category  2.411006
order_protocol          0.503981
total_items             0.000000
subtotal               0.000000
num_distinct_items      0.000000
min_item_price          0.000000
max_item_price          0.000000
total_onshift_partners  8.236927
total_busy_partners     8.236927
total_outstanding_orders 8.236927
dtype: float64
```

- most missing values are in "total_onshift_partners", "total_busy_partners", "total_outstanding_orders" have same amount of missing value i.e. 8.24%
- followed by 2nd highest is store_primary_category
- "actual_delivery_time" have lowest number of missing value (0.0035%)
- we can impute the missing values in "total_onshift_partners", "total_busy_partners", "total_outstanding_orders" using group by market_id and mean values per market.

```
In [12]: 1 for i in df:
2         print(f" unique values in {i:<30}----- => { df[i].nunique()}")
```

```
unique values in market_id          ----- => 6
unique values in created_at          ----- => 180985
unique values in actual_delivery_time ----- => 178110
unique values in store_id            ----- => 6743
unique values in store_primary_category ----- => 74
unique values in order_protocol       ----- => 7
unique values in total_items          ----- => 57
unique values in subtotal             ----- => 8368
unique values in num_distinct_items   ----- => 20
unique values in min_item_price       ----- => 2312
unique values in max_item_price       ----- => 2652
unique values in total_onshift_partners ----- => 172
unique values in total_busy_partners  ----- => 159
unique values in total_outstanding_orders ----- => 281
```

Checking "store_primary_category" feature structure and characteristic

```
In [13]: 1 df["store_primary_category"].nunique()
```

Out[13]: 74

```
In [14]: 1 df["store_primary_category"].value_counts()
```

```
Out[14]: store_primary_category
american      19399
pizza         17321
mexican       17099
burger        10958
sandwich      10060
...
lebanese       9
belgian        2
indonesian     2
chocolate      1
alcohol-plus-food 1
Name: count, Length: 74, dtype: int64
```

```
In [15]: 1 df["store_primary_category"].unique()
```

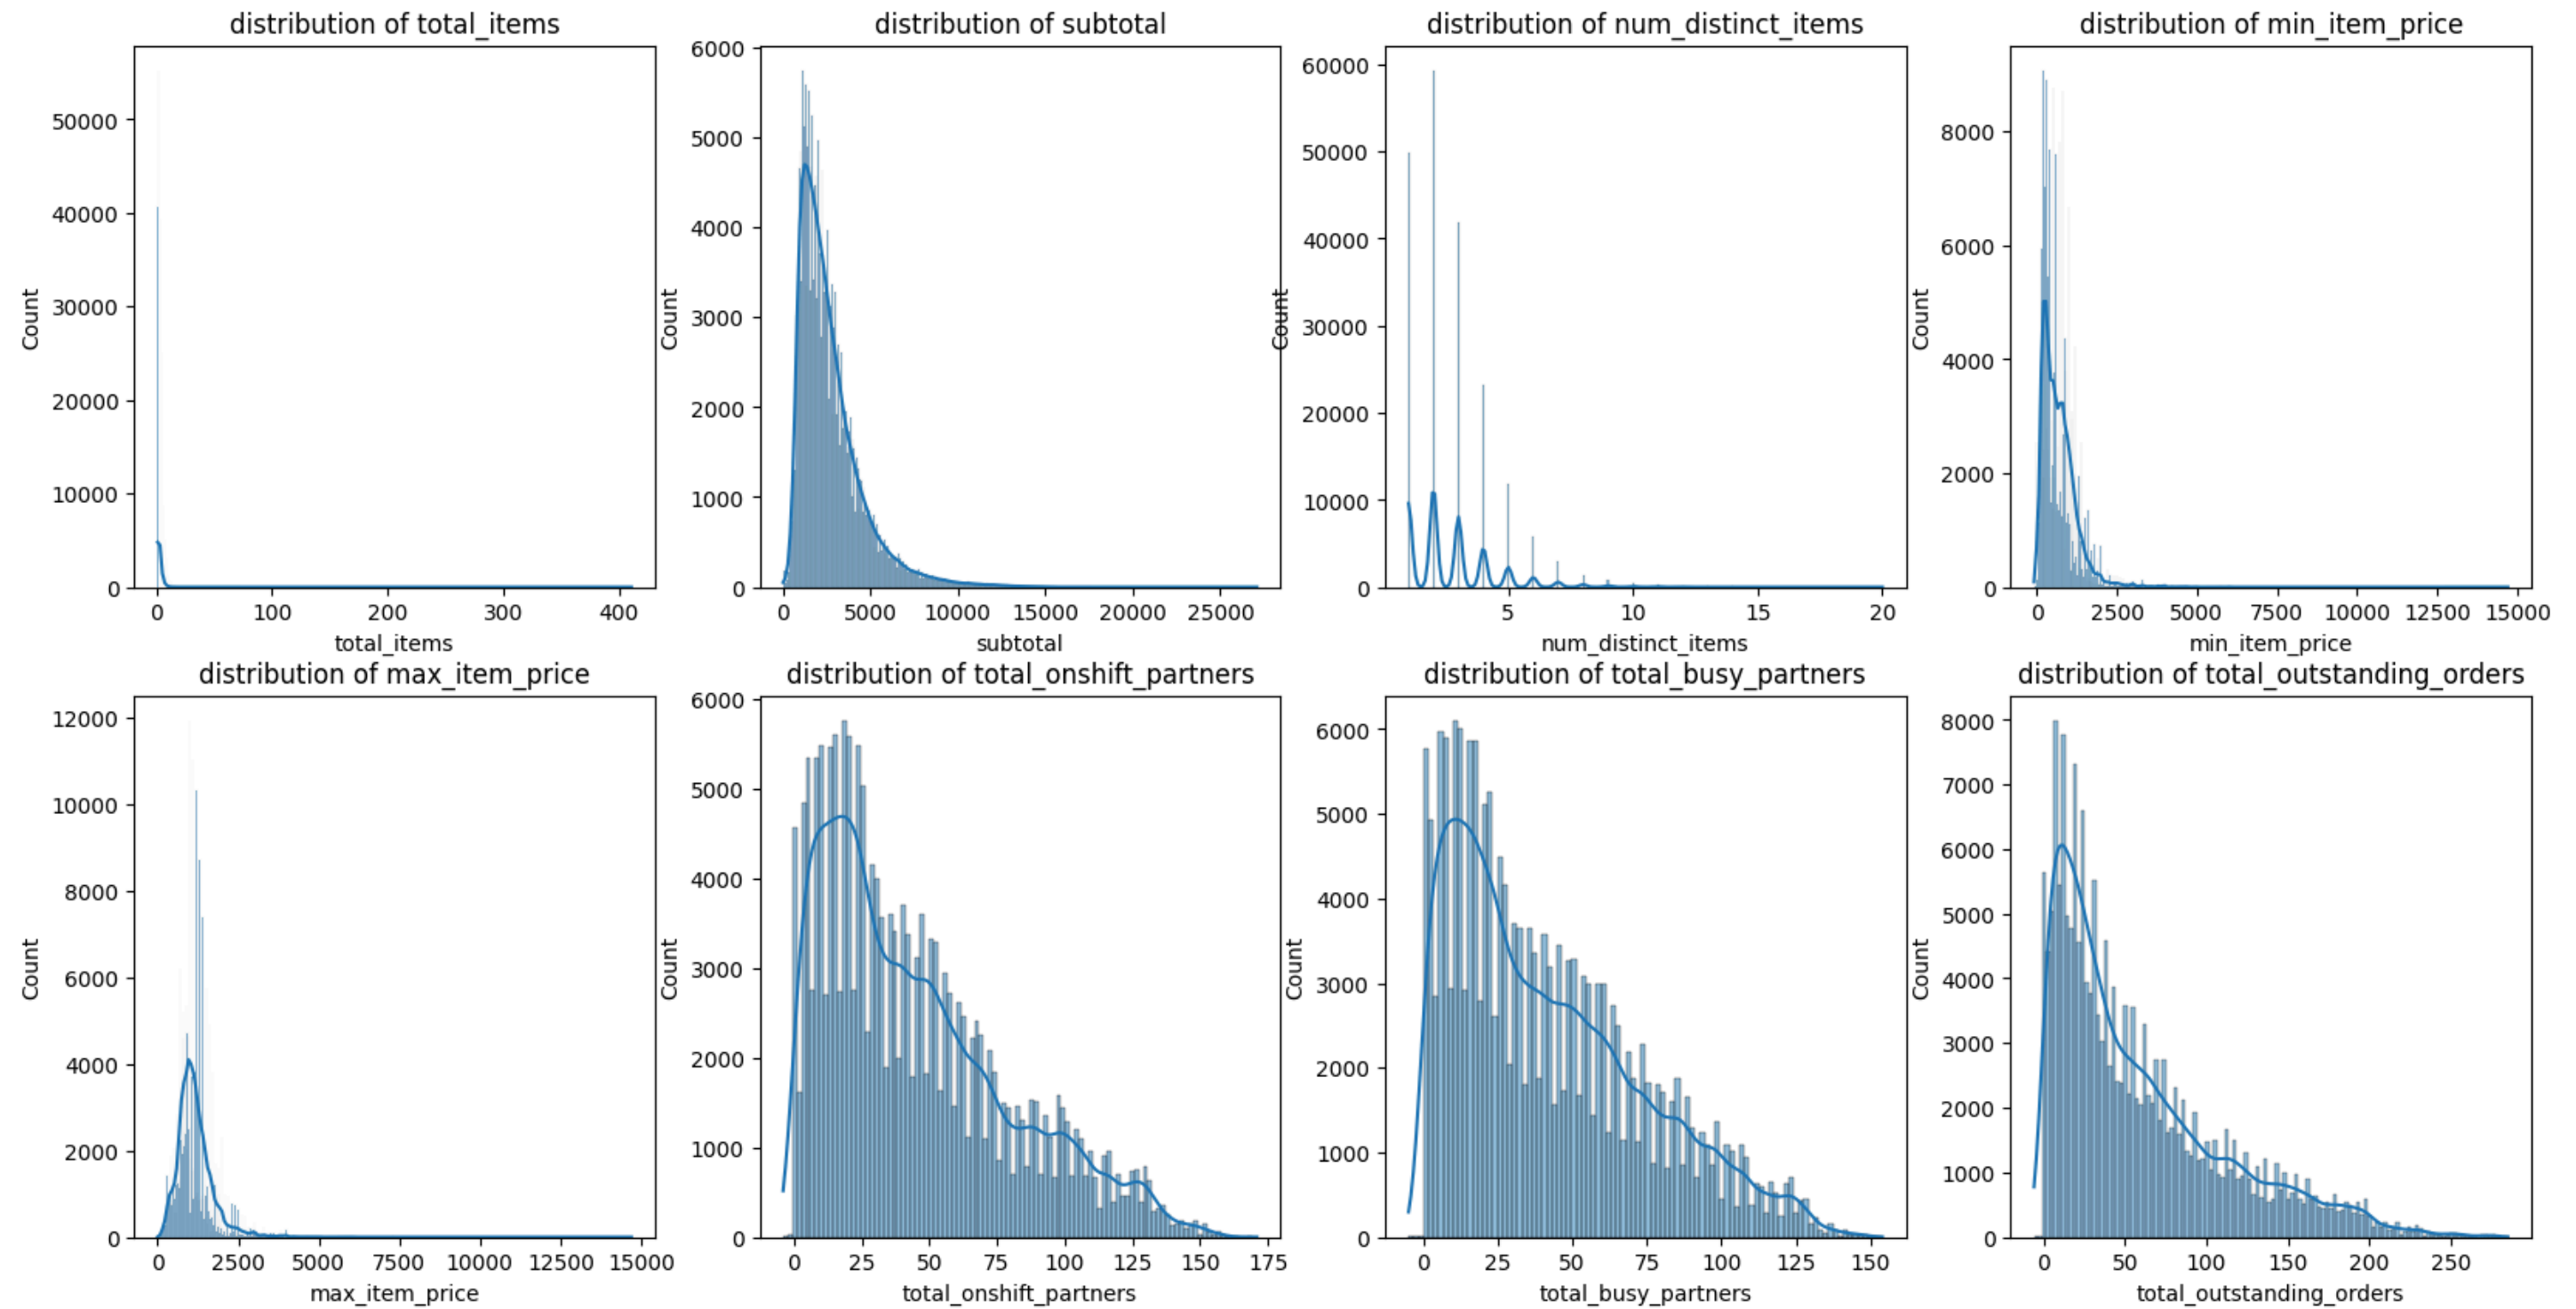
```
Out[15]: array(['american', 'mexican', nan, 'indian', 'italian', 'sandwich',
               'thai', 'cafe', 'salad', 'pizza', 'chinese', 'singaporean',
               'burger', 'breakfast', 'mediterranean', 'japanese', 'greek',
               'catering', 'filipino', 'convenience-store', 'other', 'korean',
               'vegan', 'asian', 'barbecue', 'fast', 'dessert', 'smoothie',
               'seafood', 'vietnamese', 'cajun', 'steak', 'middle-eastern',
               'soup', 'vegetarian', 'persian', 'nepalese', 'sushi',
               'latin-american', 'hawaiian', 'chocolate', 'burmese', 'british',
               'pasta', 'alcohol', 'dim-sum', 'peruvian', 'turkish', 'malaysian',
               'ethiopian', 'afghan', 'bubble-tea', 'german', 'french',
               'caribbean', 'gluten-free', 'comfort-food', 'gastropub',
               'pakistani', 'moroccan', 'spanish', 'southern', 'tapas', 'russian',
               'brazilian', 'european', 'cheese', 'african', 'argentine',
               'kosher', 'irish', 'lebanese', 'belgian', 'indonesian',
               'alcohol-plus-food'], dtype=object)
```

Checking distribution of all the numerical feature

```
In [16]: 1 numerical_columns = df.select_dtypes(include=['int', 'float']).columns.tolist()
          2 numerical_columns= numerical_columns[2:]
          3 numerical_columns
```

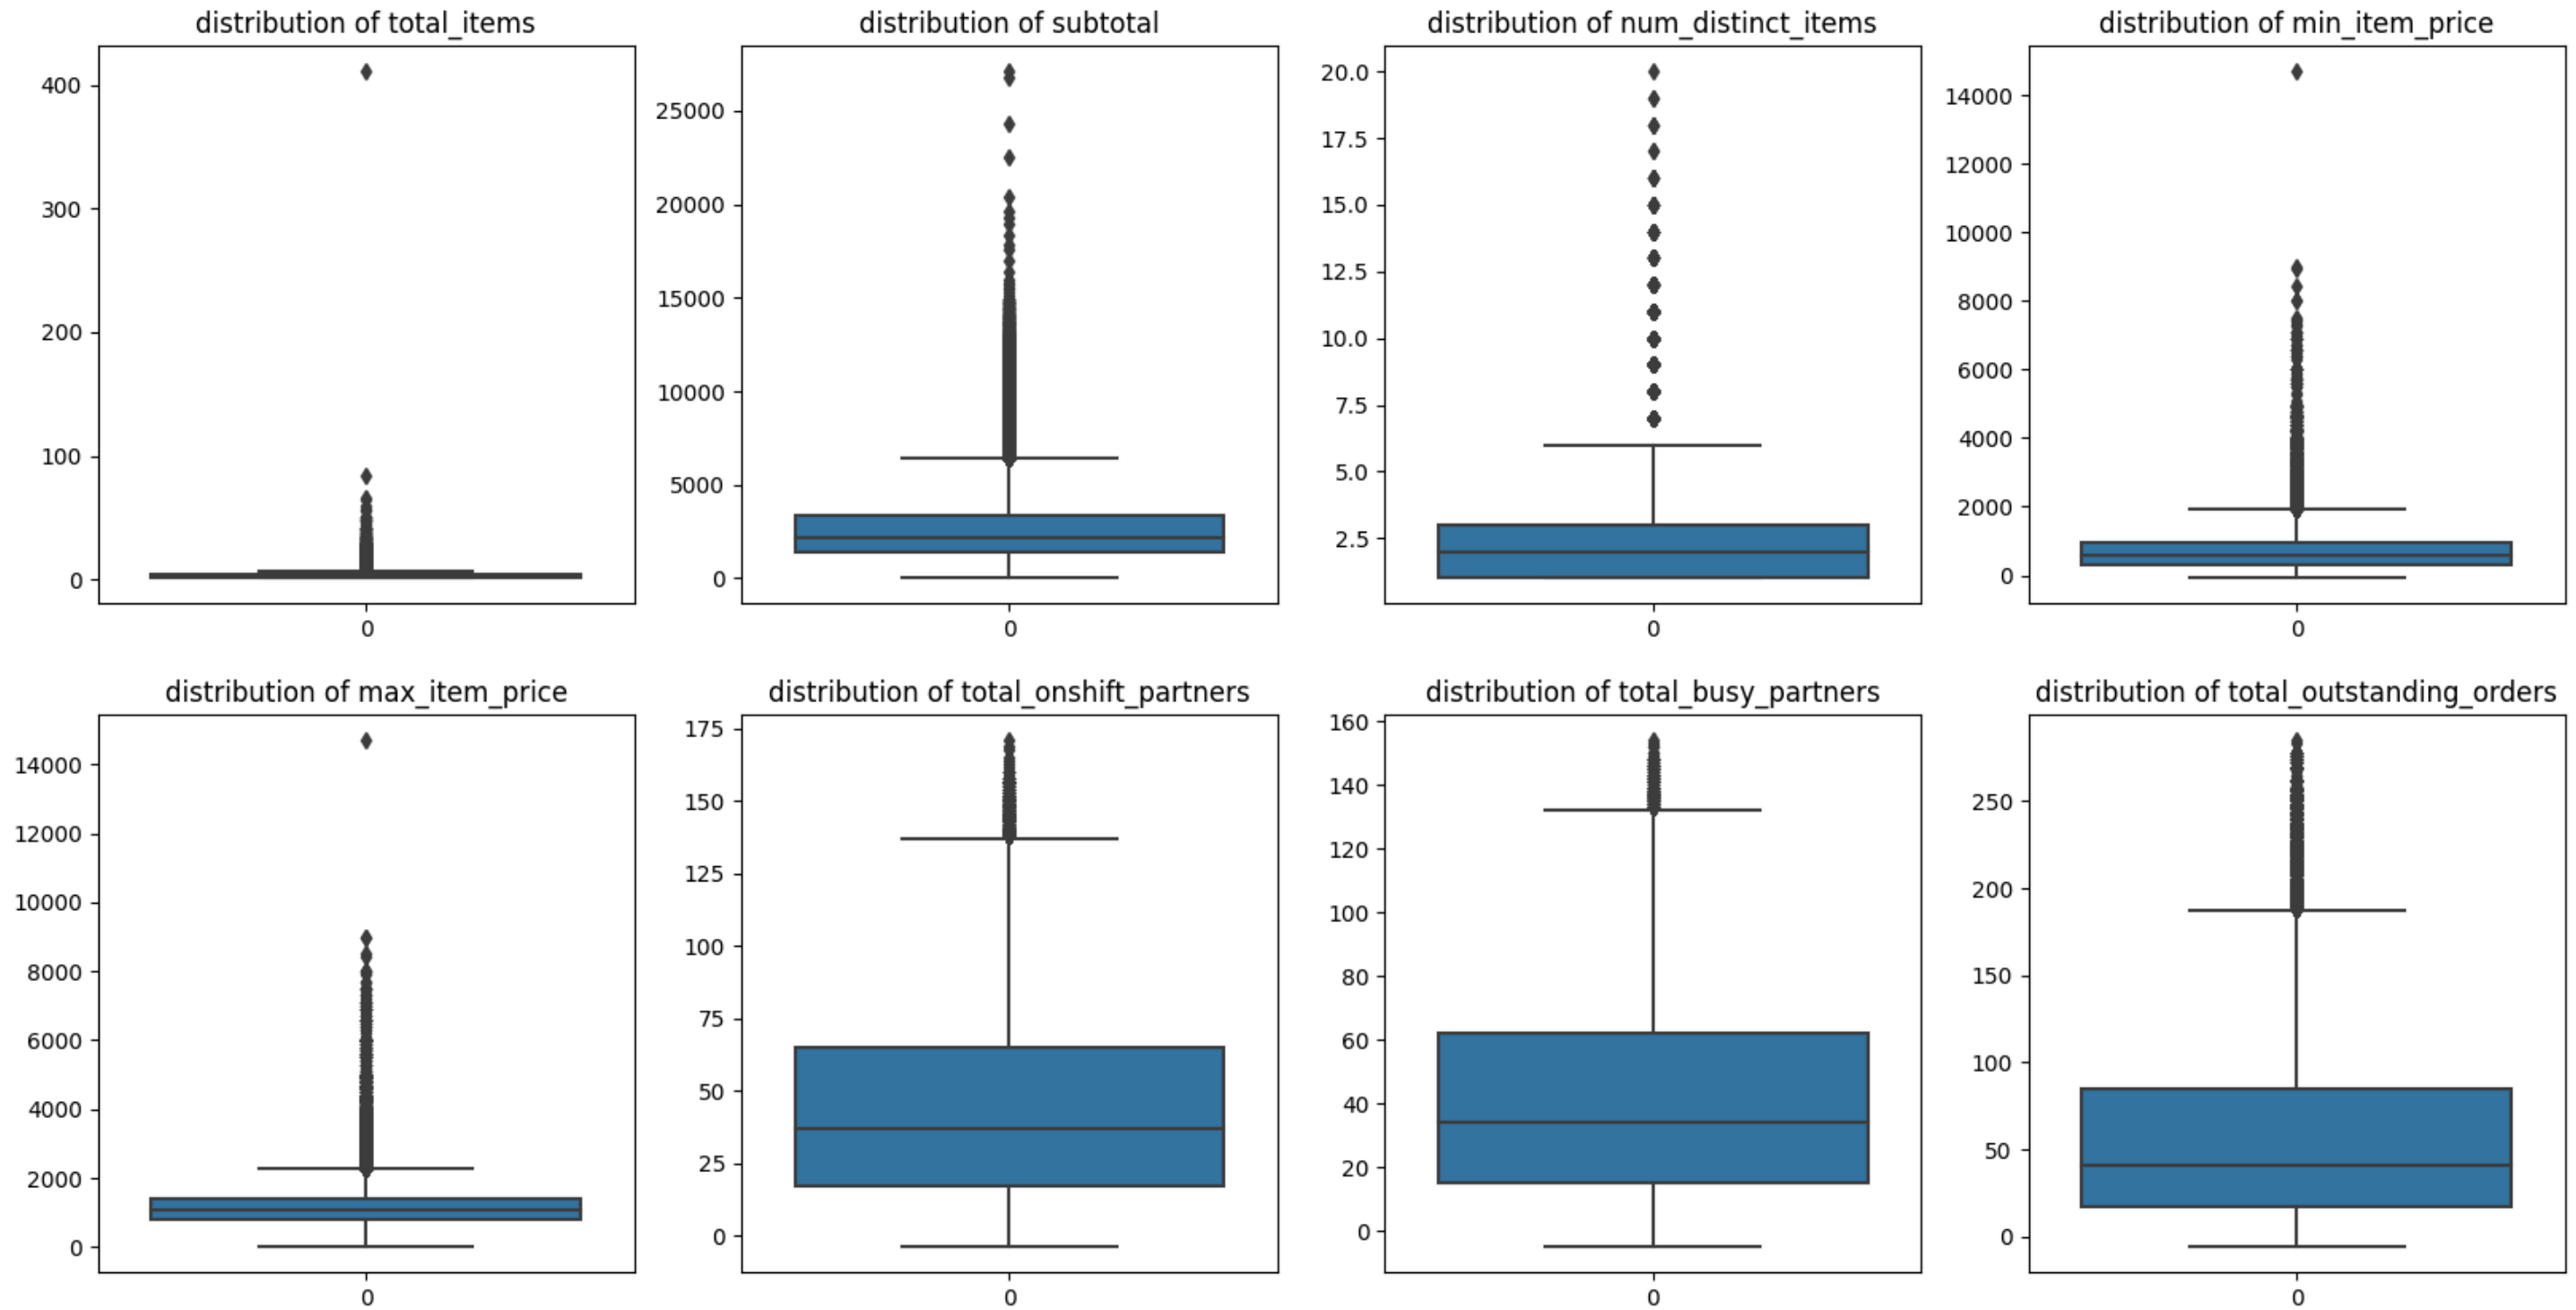
```
Out[16]: ['total_items',
          'subtotal',
          'num_distinct_items',
          'min_item_price',
          'max_item_price',
          'total_onshift_partners',
          'total_busy_partners',
          'total_outstanding_orders']
```

```
In [17]: 1 fig = plt.figure(figsize=(20, 10))
2
3 for i in range(len(numerical_columns)):
4     plt.subplot(2, 4, i+1)
5     sns.histplot(df[numerical_columns[i]], kde=True)
6     plt.title(f"distribution of {numerical_columns[i]}")
```



- All the distribution showing right skweness, indicating outliers.

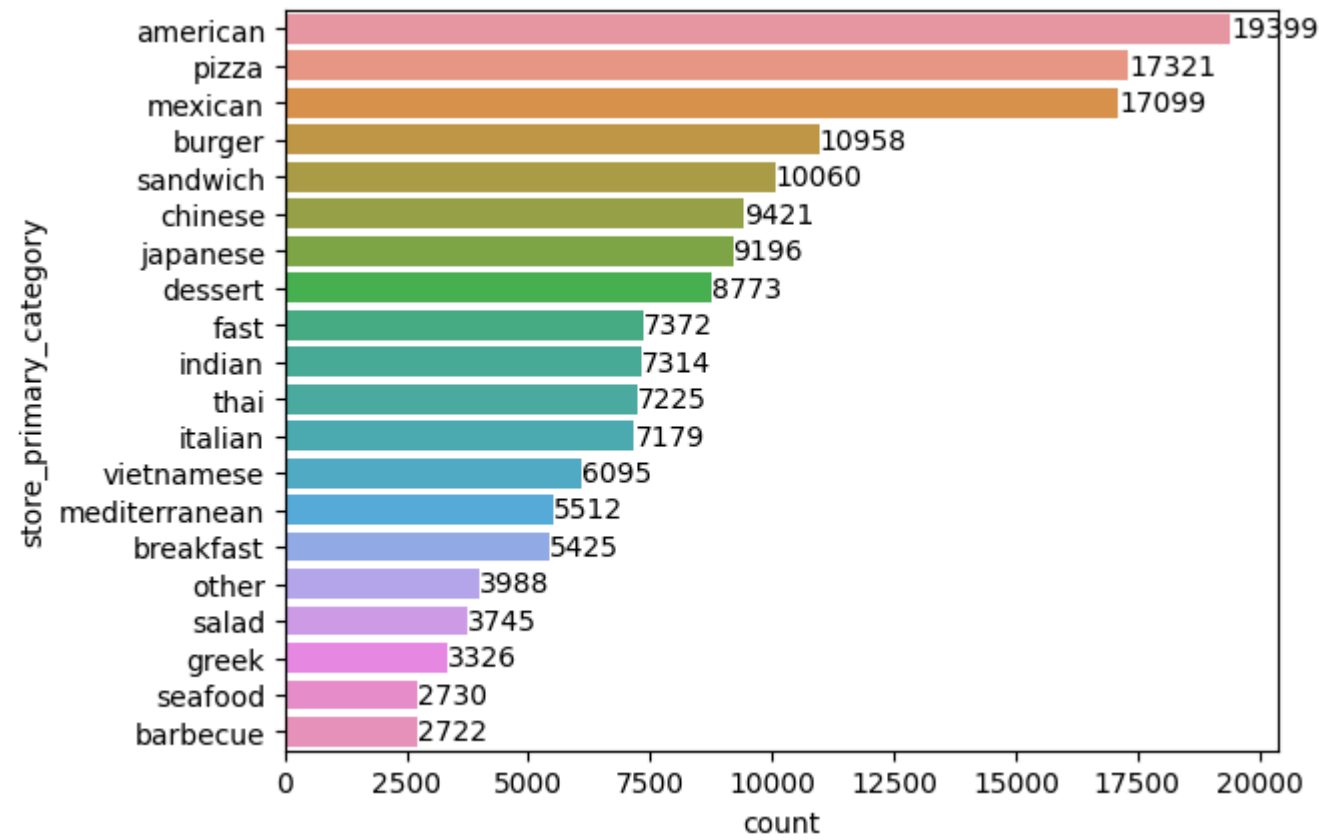
```
In [18]: 1 fig = plt.figure(figsize=(20, 10))
2
3 for i in range(len(numerical_columns)):
4     plt.subplot(2, 4, i+1)
5     sns.boxplot( df[numerical_columns[i]])
6     plt.title(f"distribution of {numerical_columns[i]}")
```



- There are some extreme outliers which maybe not following natural pattern.

Top 20 category for the restaurant based on Number of delivery from that store

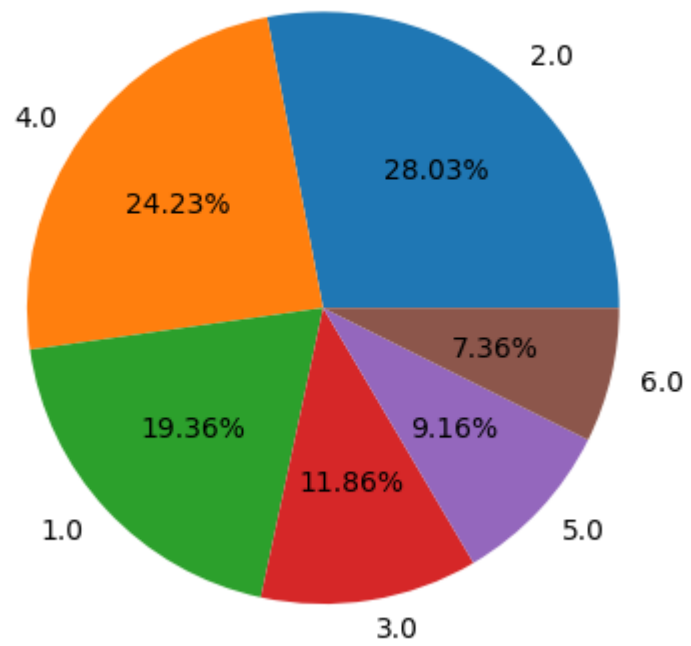

```
In [19]: 1 lbl= sns.barplot(x= df["store_primary_category"].value_counts()[ :20], y= df["store_primary_category"].value_counts()[ :20].index)
2         for i in lbl.containers:
3             lbl.bar_label(i)
4
5         plt.show()
```



```
In [20]: 1 df["market_id"].value_counts()
```

```
Out[20]: market_id
2.0    55058
4.0    47599
1.0    38037
3.0    23297
5.0    18000
6.0    14450
Name: count, dtype: int64
```

```
In [21]: 1 plt.pie(df["market_id"].value_counts(), labels= df["market_id"].value_counts().index, autopct='%.2f%%' )
2 plt.show()
```



- 28% order are from Market_id 2.0 followed by market_id 4.0
- Market_id 6.0 has lowest orders.

- **We have created_at and actual_delivery_time, from this we can calculate total time taken to deliver i.e. dependent variable.**
- By using timedelta we can convert time_taken into minutes.

```
In [22]: 1 df["time_taken"] = df["actual_delivery_time"] - df["created_at"]
```

```
In [23]: 1 df["time_taken_min"] = pd.to_timedelta(df["time_taken"])/pd.Timedelta("60s")
```

In [24]:

1df.head()

Out[24]:

| | market_id | created_at | actual_delivery_time | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_pa |
|---|-----------|---------------------|----------------------|----------------------------------|------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|---------------|
| 0 | 1.0 | 2015-02-06 22:24:17 | 2015-02-06 23:27:16 | df263d996281d984952c07998dc54358 | american | 1.0 | 4 | 3441 | 4 | 557 | 1239 | 33.0 | |
| 1 | 2.0 | 2015-02-10 21:49:25 | 2015-02-10 22:56:29 | f0ade77b43923b38237db569b016ba25 | mexican | 2.0 | 1 | 1900 | 1 | 1400 | 1400 | 1.0 | |
| 2 | 3.0 | 2015-01-22 20:39:28 | 2015-01-22 21:09:09 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | 1 | 1900 | 1 | 1900 | 1900 | 1.0 | |
| 3 | 3.0 | 2015-02-03 21:21:45 | 2015-02-03 22:13:00 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | 6 | 6900 | 5 | 600 | 1800 | 1.0 | |
| 4 | 3.0 | 2015-02-15 02:40:36 | 2015-02-15 03:20:26 | f0ade77b43923b38237db569b016ba25 | NaN | 1.0 | 3 | 3900 | 3 | 1100 | 1600 | 6.0 | |

- Deriving the feature hour_created_at and day of week at which order was created from created_at

In [25]:

1df["hour_created"]= df["created_at"].dt.hour
2df["day_of_week"]= df["created_at"].dt.dayofweek

checking distribution of target column i.e. time_taken_min

In [26]:

1df["time_taken_min"].describe()

Out[26]:

```
count    197421.000000  
mean       48.470956  
std       320.493482  
min        1.683333  
25%       35.066667  
50%       44.333333  
75%       56.350000  
max      141947.650000  
Name: time_taken_min, dtype: float64
```

- 75% of the time time taken to deliver is less than or equal to 56 minutes.
- mean time to deliver the order is 48.47 minutes.

checking for orders with more than 1000 minutes

In [27]:

1df.loc[df["time_taken_min"]>1000]

Out[27]:

| | market_id | created_at | actual_delivery_time | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy |
|--------|-----------|---------------------|----------------------|----------------------------------|------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|------------|
| 2690 | 1.0 | 2014-10-19 05:24:15 | 2015-01-25 19:11:54 | 675f9820626f5bc0afb47b57890b466e | italian | 1.0 | 1 | 1695 | 1 | 1595 | 1595 | NaN | |
| 27189 | 1.0 | 2015-02-16 02:24:09 | 2015-02-19 22:45:31 | d397c2b2be2178fe6247bd50fc97cff2 | indian | 3.0 | 4 | 4980 | 4 | 995 | 1795 | 72.0 | |
| 185550 | 4.0 | 2015-01-28 08:34:06 | 2015-02-01 16:25:25 | 1679091c5a880faf6fb5e6087eb1b2dc | dessert | 5.0 | 3 | 1520 | 3 | 220 | 750 | 0.0 | |

- There are only 3 orders which took more than 5000 minutes for delivery.
- we can drop this

In [28]:

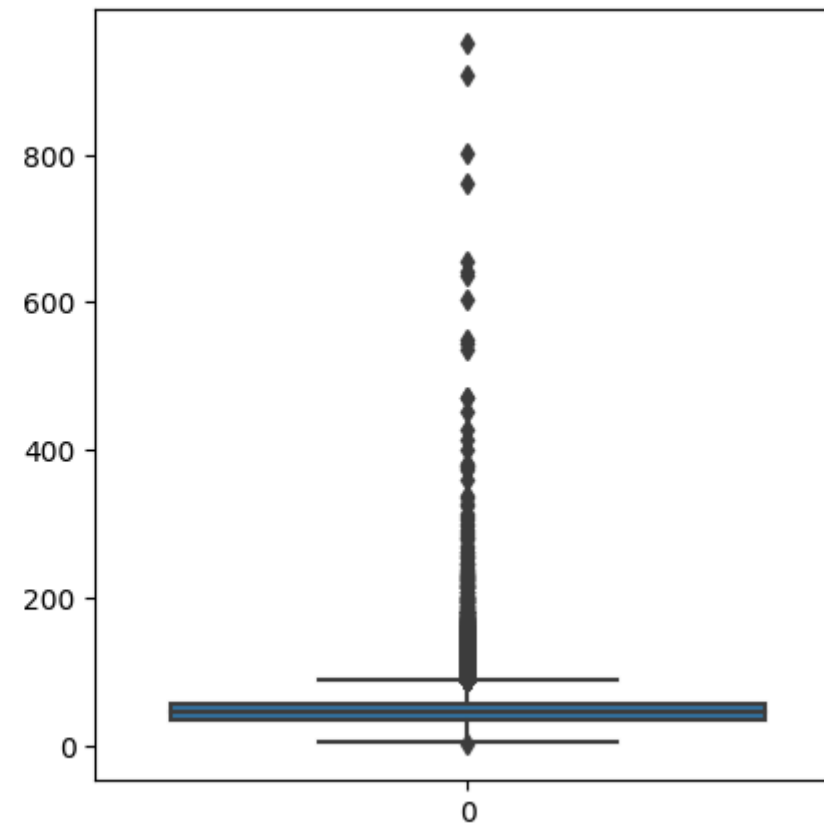
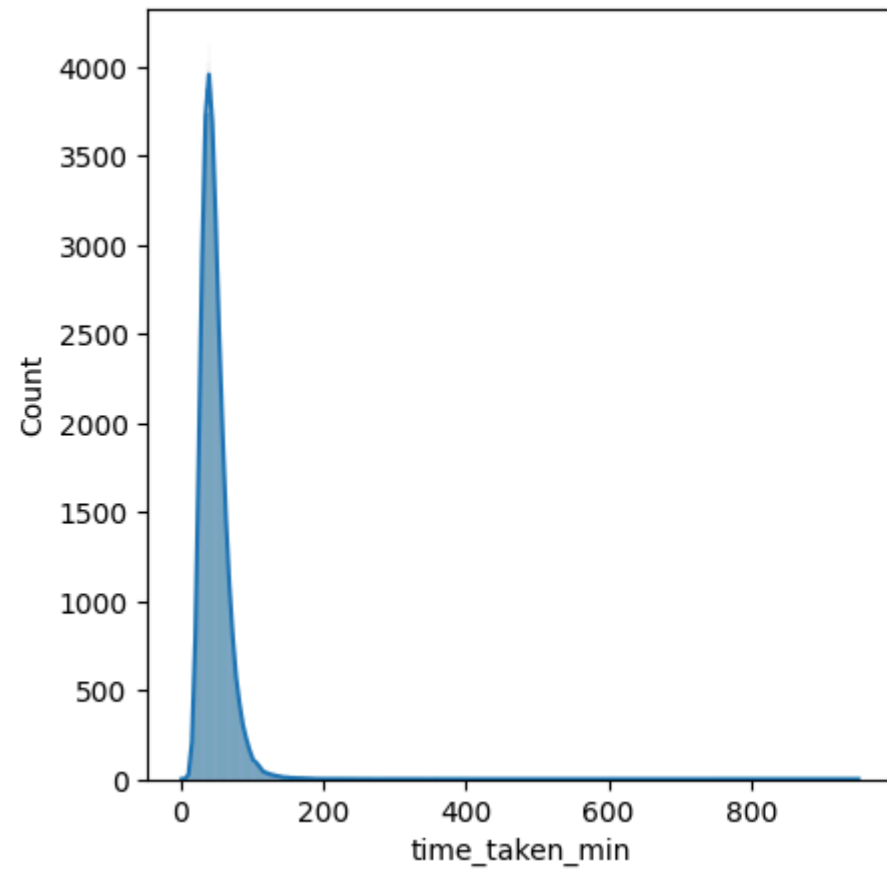
1df= df.loc[df["time_taken_min"]<1000]

In [29]:

1df.shape

Out[29]: (197418, 18)

```
In [30]: 1 plt.figure(figsize = (12, 5))
2         plt.subplot(1, 2, 1)
3         plt.subplots_adjust(wspace= 0.5)
4         sns.histplot(df['time_taken_min'], kde=True)
5
6
7         plt.subplot(1, 2, 2)
8         sns.boxplot(df['time_taken_min'])
9
10        plt.show()
```



Top stores with most orders in percentage

```
In [31]: 1 df["store_id"].value_counts(normalize=True)*100
```

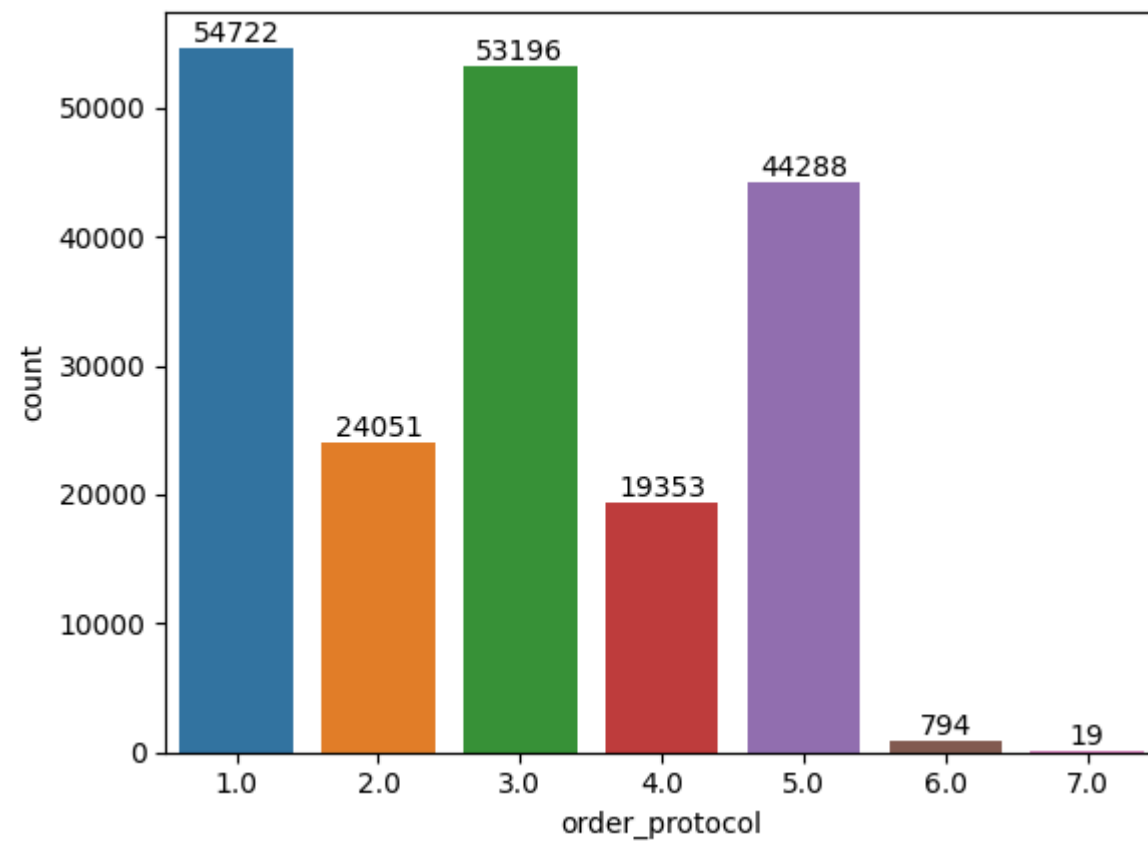
```
Out[31]: store_id
d43ab110ab2489d6b9b2caa394bf920f    0.474627
757b505cfd34c64c85ca5b5690ee5293    0.437144
faacbcd5bf1d018912c116bf2783e9a1    0.412323
cfecdb276f634854f3ef915e2e980c31    0.387503
45c48cce2e2d7fbdea1afc51c7c6ad26    0.365215
...
adad0f2b196a1ed3e3b9d9025c397132    0.000507
2e6d9c6052e99fcdfa61d9b9da273ca2    0.000507
25daeb9b3072e9c53f66a2196a92a011    0.000507
55285adfd78a019a3245917649e29b3c    0.000507
df263d996281d984952c07998dc54358    0.000507
Name: proportion, Length: 6743, dtype: float64
```

Order Protocol- how the order was placed.

```
In [32]: 1 # In percentage
        2 df["order_protocol"].value_counts(normalize=True)*100
```

```
Out[32]: order_protocol
1.0      27.859263
3.0      27.082368
5.0      22.547258
2.0      12.244493
4.0       9.852716
6.0       0.404230
7.0       0.009673
Name: proportion, dtype: float64
```

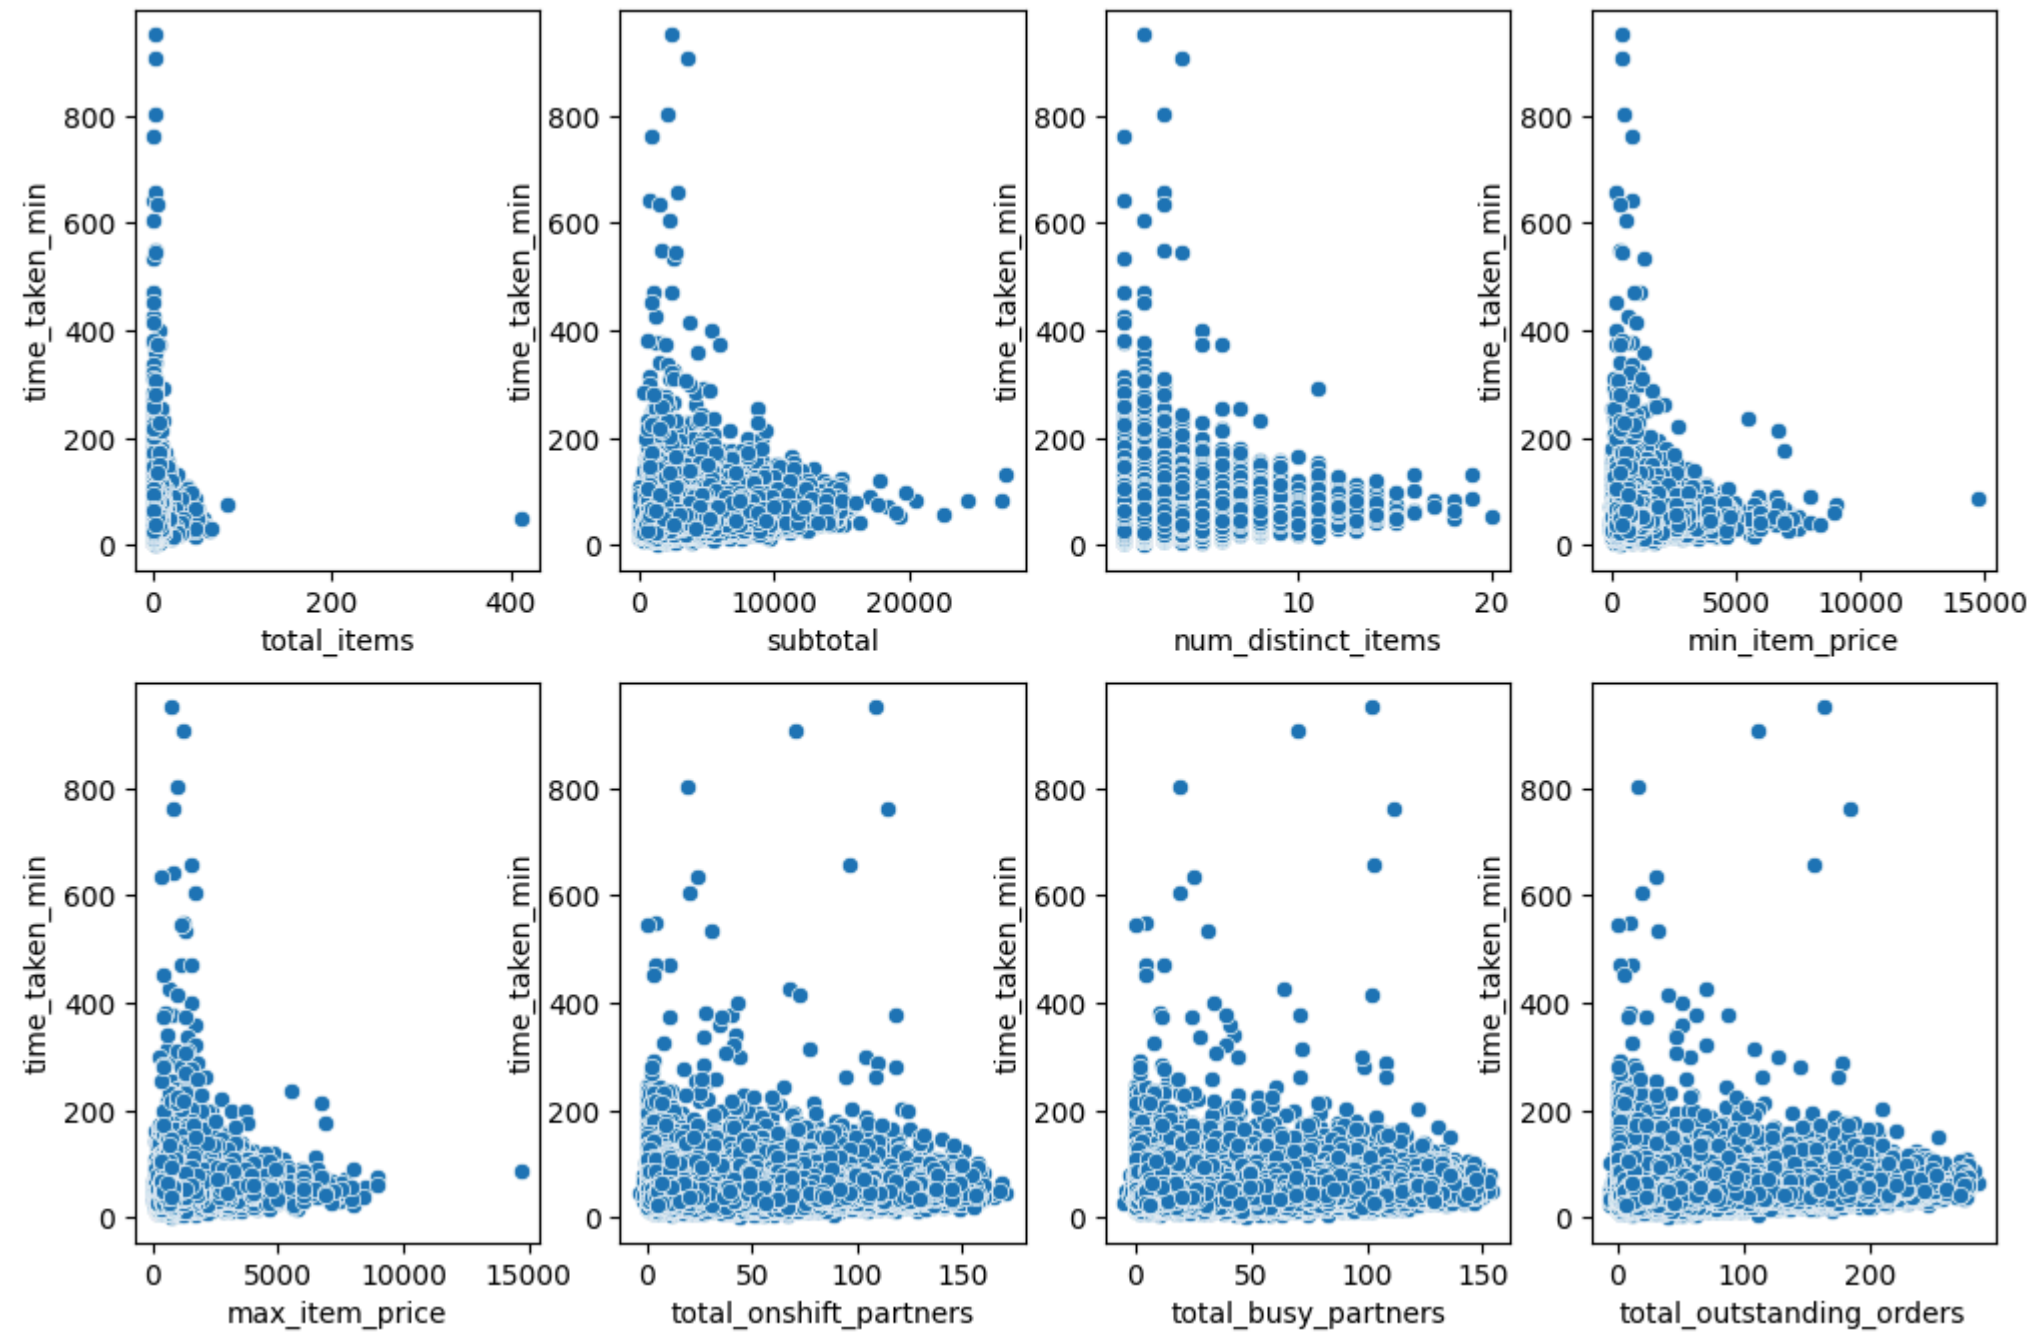
```
In [33]: 1 labl= sns.countplot(data=df, x=df["order_protocol"])
        2 for i in labl.containers:
        3     labl.bar_label(i)
        4 plt.show()
```



- Most of the order was placed through "1.0" and "3.0" type of mode.
- Very Few order was place through "7.0"

In [34]:

```
1 rows, cols = 2,4
2 fig, axs = plt.subplots(rows,cols, figsize=(12, 8))
3 index = 0
4 for row in range(rows):
5     for col in range(cols):
6         sns.scatterplot(x=numerical_columns[index], y=df["time_taken_min"], data=df, ax=axs[row, col])
7         index += 1
8
```



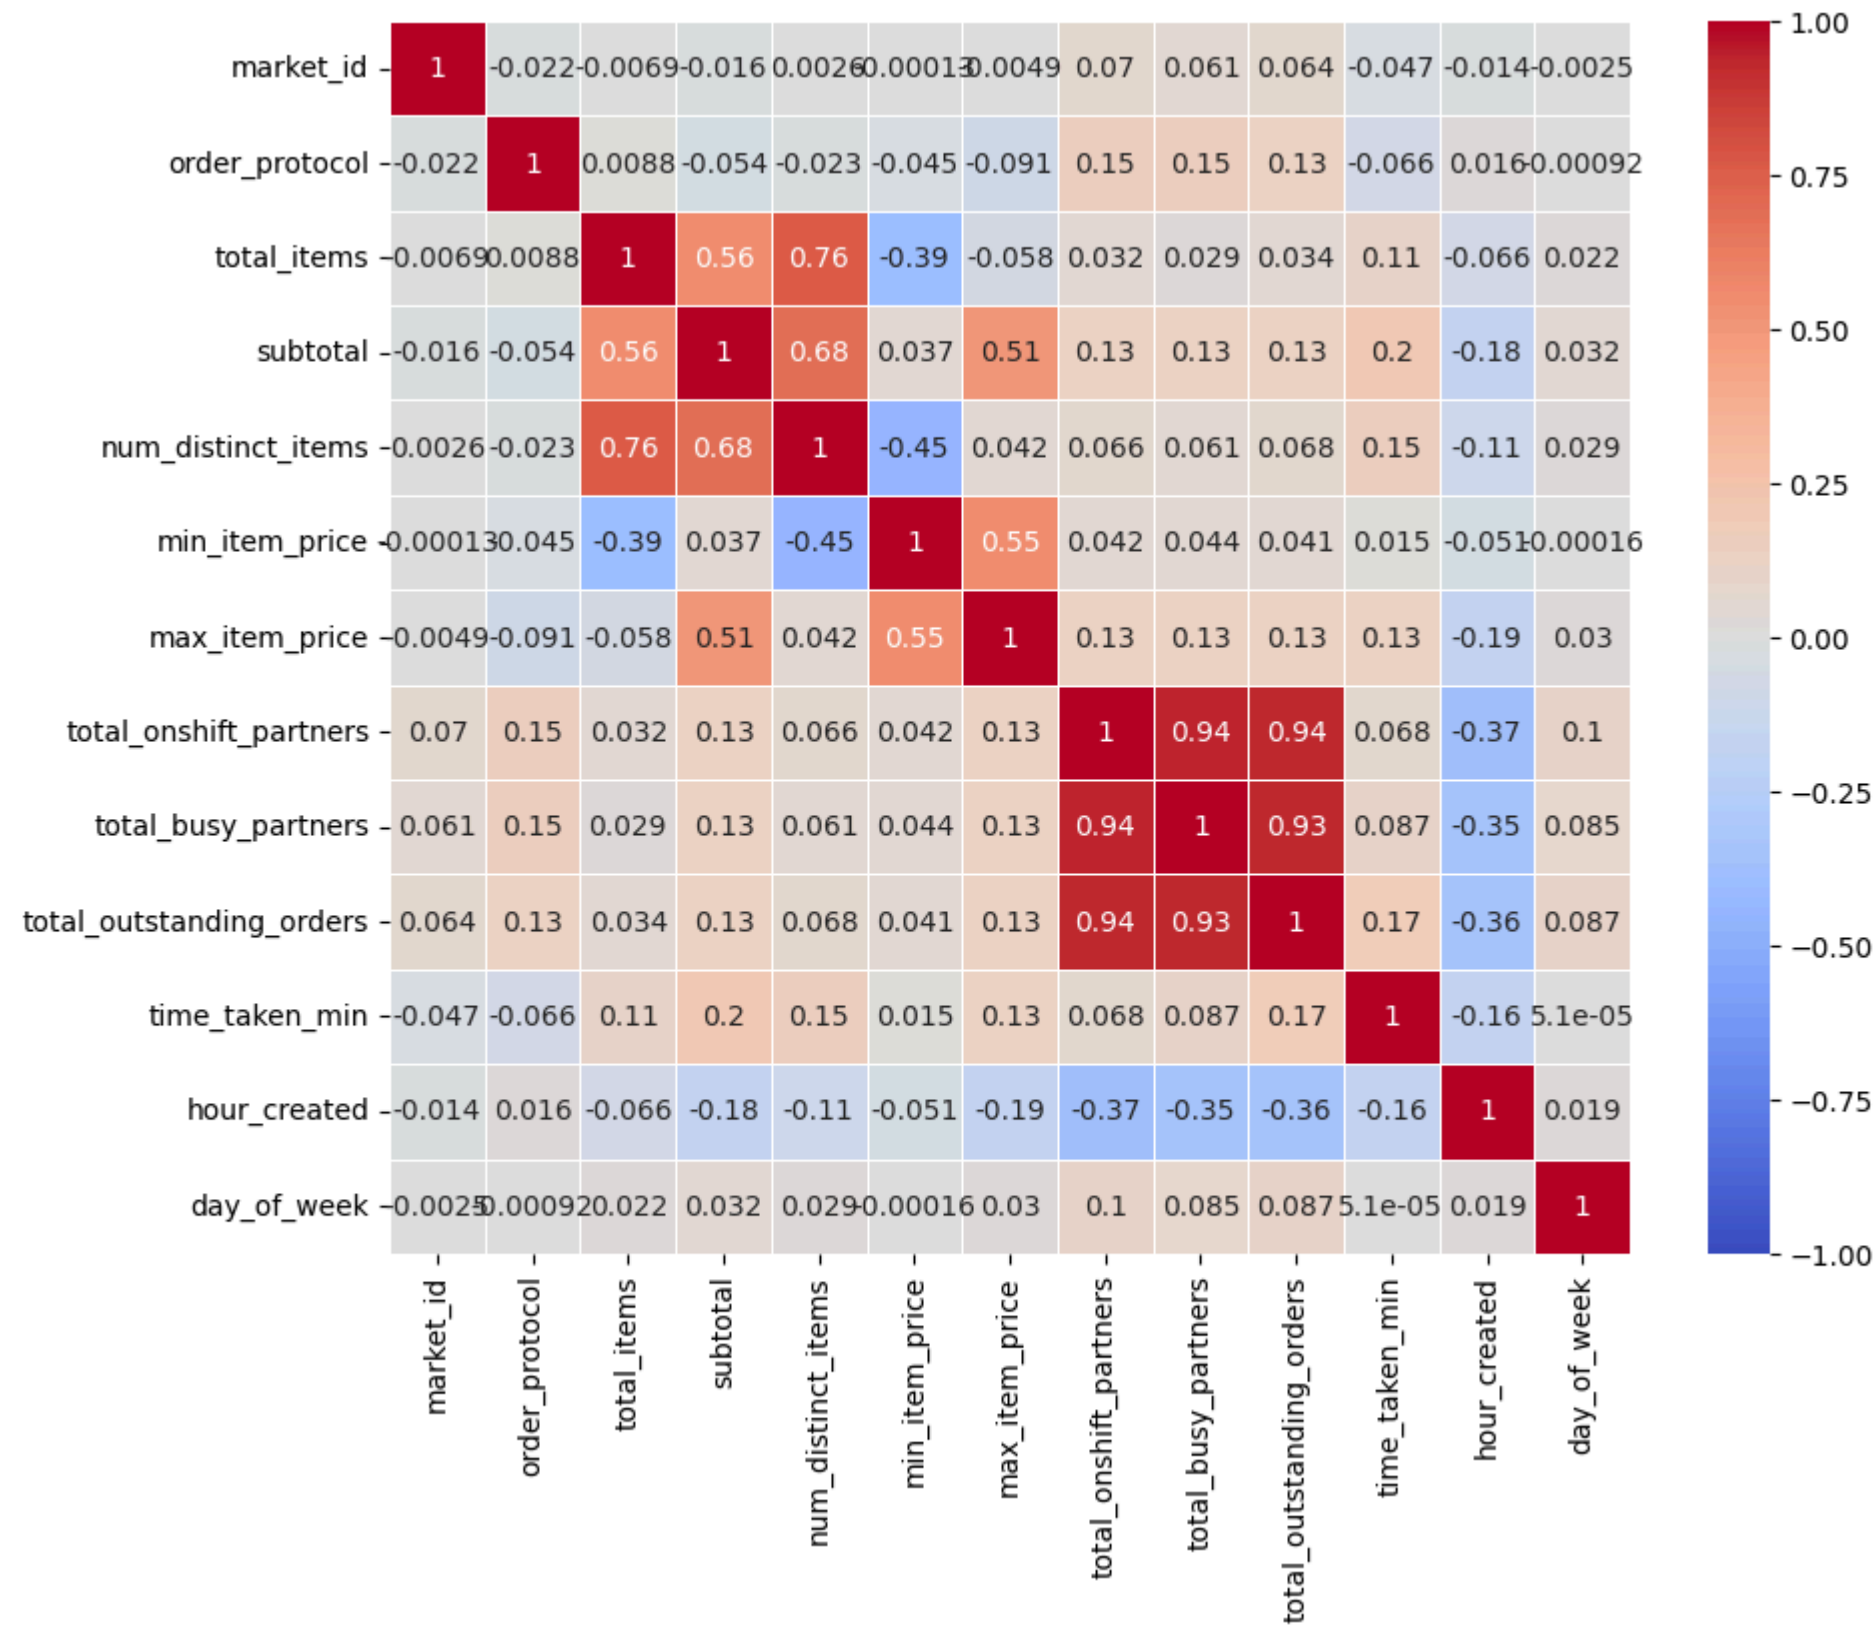
- As we can see there is not much impact on delivery of order with different variable like total_items, max_item_price, busy_partners.

```
In [35]: 1 numerical_columns = df.select_dtypes(include=['int', 'float']).columns.tolist()
2         #numerical_columns= numerical_columns[2:]
3         numerical_columns
```

```
Out[35]: ['market_id',
'order_protocol',
'total_items',
'subtotal',
'num_distinct_items',
'min_item_price',
'max_item_price',
'total_onshift_partners',
'total_busy_partners',
'total_outstanding_orders',
'time_taken_min',
'hour_created',
'day_of_week']
```


In [36]:

```
1 plt.figure(figsize = (10,8))
2 dt_corr= df[numerical_columns].corr()
3 sns.heatmap(dt_corr, annot= True, cmap='coolwarm', linewidth=.5, vmin=-1, vmax=1)
4 plt.show()
```



Group by market to get min, max and Avg time taken market wise

```
In [37]: 1 df.groupby(["market_id"])["time_taken_min"].agg({"min", "max", "mean"}).reset_index().sort_values(by = "mean", ascending=False)
```

Out[37]:

| | market_id | max | mean | min |
|---|-----------|------------|-----------|-----------|
| 0 | 1.0 | 907.450000 | 51.331649 | 6.766667 |
| 2 | 3.0 | 802.966667 | 47.695185 | 7.833333 |
| 5 | 6.0 | 640.950000 | 47.258626 | 9.283333 |
| 3 | 4.0 | 535.983333 | 47.236536 | 6.433333 |
| 4 | 5.0 | 381.866667 | 46.495435 | 10.316667 |
| 1 | 2.0 | 950.533333 | 46.073661 | 1.683333 |

```
In [38]: 1 #cheking duplicate record
2 df.duplicated().sum()
```

Out[38]: 0

Treating missing values

```
In [39]: 1 df.isna().sum()/len(df)*100
```

Out[39]:

| | |
|--------------------------|----------|
| market_id | 0.499954 |
| created_at | 0.000000 |
| actual_delivery_time | 0.000000 |
| store_id | 0.000000 |
| store_primary_category | 2.411128 |
| order_protocol | 0.504007 |
| total_items | 0.000000 |
| subtotal | 0.000000 |
| num_distinct_items | 0.000000 |
| min_item_price | 0.000000 |
| max_item_price | 0.000000 |
| total_onshift_partners | 8.236838 |
| total_busy_partners | 8.236838 |
| total_outstanding_orders | 8.236838 |
| time_taken | 0.000000 |
| time_taken_min | 0.000000 |
| hour_created | 0.000000 |
| day_of_week | 0.000000 |

dtype: float64

```
In [40]: 1 df['total_onshift_partners'] = df.groupby('market_id')['total_onshift_partners'].transform(lambda x: x.fillna(x.mean()))
2 df['total_busy_partners'] = df.groupby('market_id')['total_busy_partners'].transform(lambda x: x.fillna(x.mean()))
3 df['total_outstanding_orders'] = df.groupby('market_id')['total_outstanding_orders'].transform(lambda x: x.fillna(x.mean()))
4
5
```

```
In [41]: 1 def fill_mode(col):
2         mode_value = col.mode()
3         if len(mode_value) > 0:
4             return col.fillna(mode_value.iloc[0])
5         else:
6             return col
```

```
In [42]: 1 df['store_primary_category'] = df.groupby('store_id')['store_primary_category'].transform(fill_mode)
```

```
In [43]: 1 df.isna().sum()/len(df)*100
```

```
Out[43]: market_id          0.499954
created_at          0.000000
actual_delivery_time 0.000000
store_id            0.000000
store_primary_category 0.439170
order_protocol      0.504007
total_items         0.000000
subtotal            0.000000
num_distinct_items  0.000000
min_item_price      0.000000
max_item_price      0.000000
total_onshift_partners 0.499954
total_busy_partners  0.499954
total_outstanding_orders 0.499954
time_taken          0.000000
time_taken_min      0.000000
hour_created        0.000000
day_of_week         0.000000
dtype: float64
```

- Now we can drop null values, all the null values is less than 0.5%.

```
In [44]: 1 df = df.dropna()
```

```
In [45]: 1 df.isna().sum()/len(df)*100
```

```
Out[45]: market_id          0.0
created_at          0.0
actual_delivery_time 0.0
store_id            0.0
store_primary_category 0.0
order_protocol      0.0
total_items         0.0
subtotal            0.0
num_distinct_items  0.0
min_item_price      0.0
max_item_price      0.0
total_onshift_partners 0.0
total_busy_partners  0.0
total_outstanding_orders 0.0
time_taken          0.0
time_taken_min      0.0
hour_created        0.0
day_of_week         0.0
dtype: float64
```

- Since we calculated time_taken_min we can drop related columns

```
In [46]: 1 df.drop(['created_at', 'actual_delivery_time', 'time_taken'], axis=1, inplace=True)
```

```
In [47]: 1 df.tail(2)
```

Out[47]:

| | market_id | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_partners | total_outstanding_ord |
|--------|-----------|----------------------------------|------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|---------------------|-----------------------|
| 197426 | 1.0 | c81e155d85dae5430a8cee6f2242e82c | sandwich | 1.0 | 1 | 1175 | 1 | 535 | 535 | 7.0 | 7.0 | 1 |
| 197427 | 1.0 | c81e155d85dae5430a8cee6f2242e82c | sandwich | 1.0 | 4 | 2605 | 4 | 425 | 750 | 20.0 | 20.0 | 2 |

- feature like market_id, order_protocol should be categorical

```
In [48]: 1 cat_col= ["market_id", "order_protocol"]
2 df[cat_col]= df[cat_col].astype("str")
```

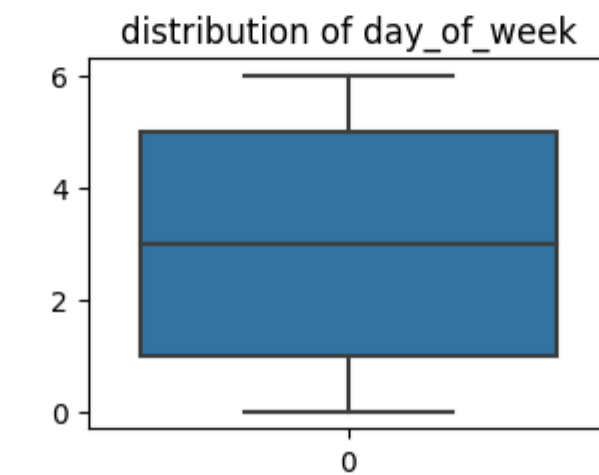
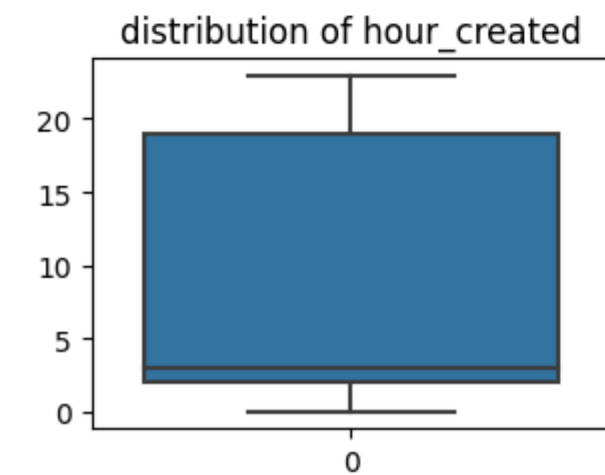
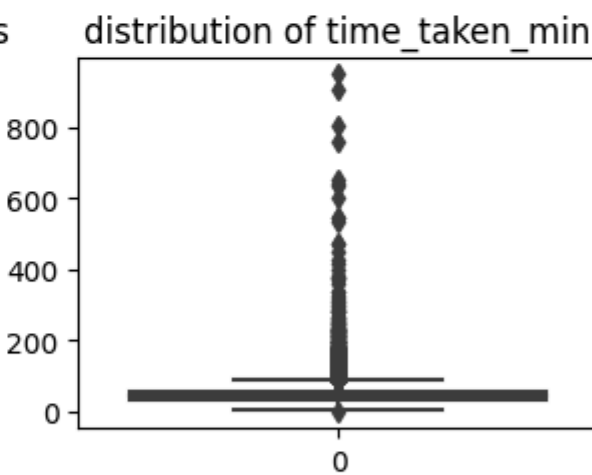
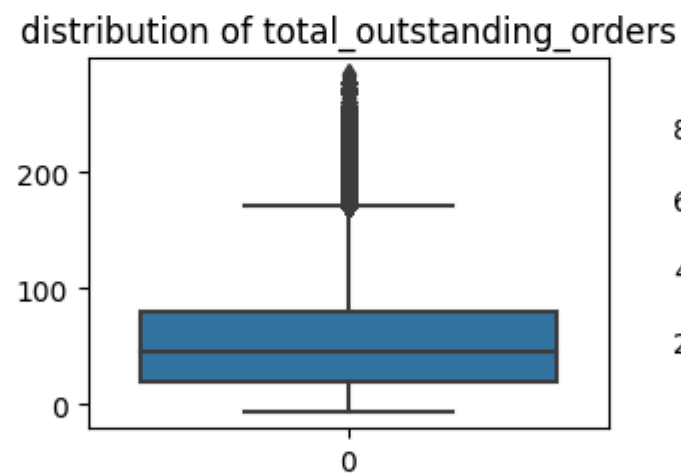
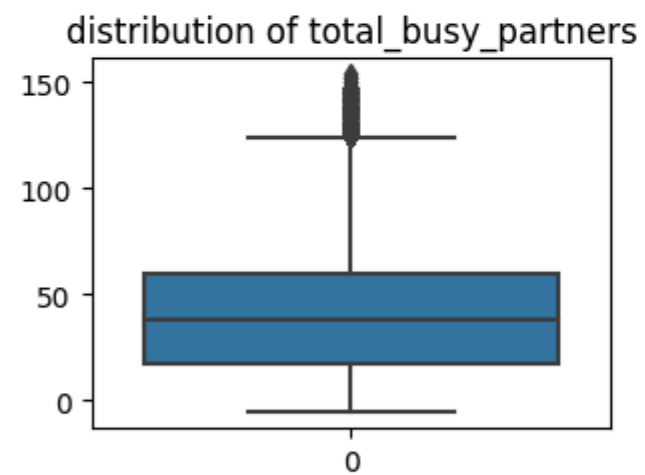
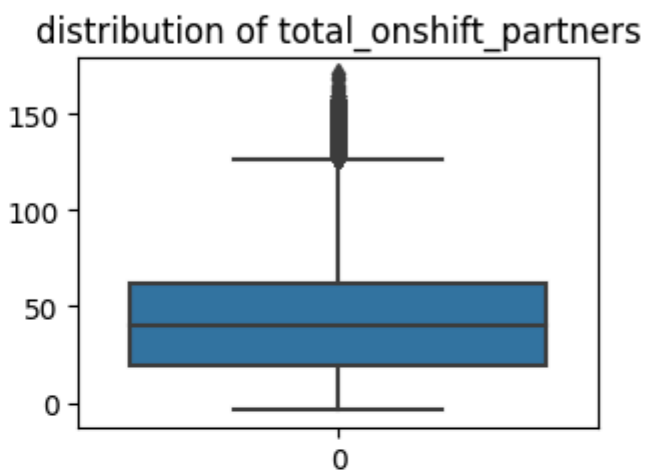
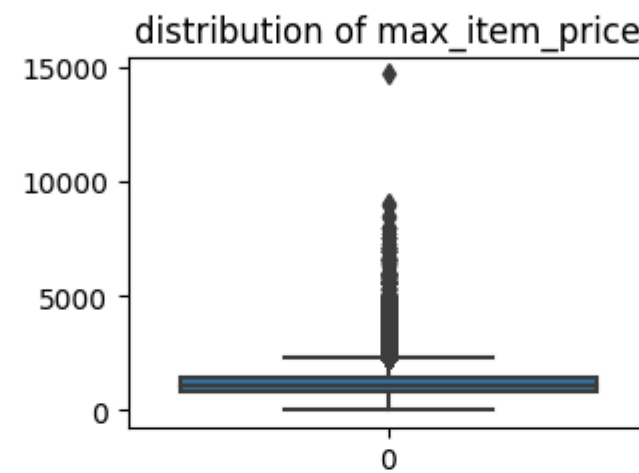
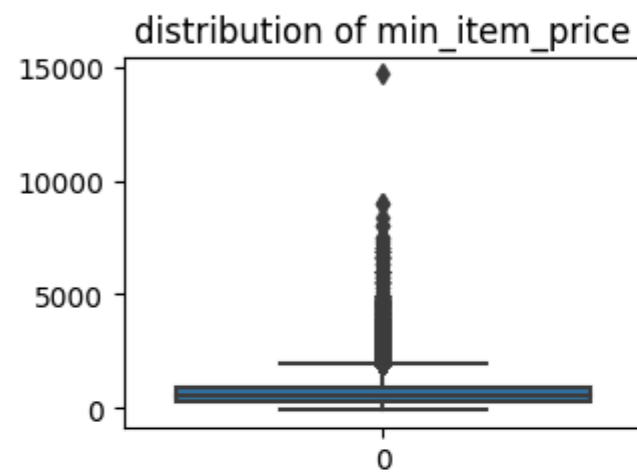
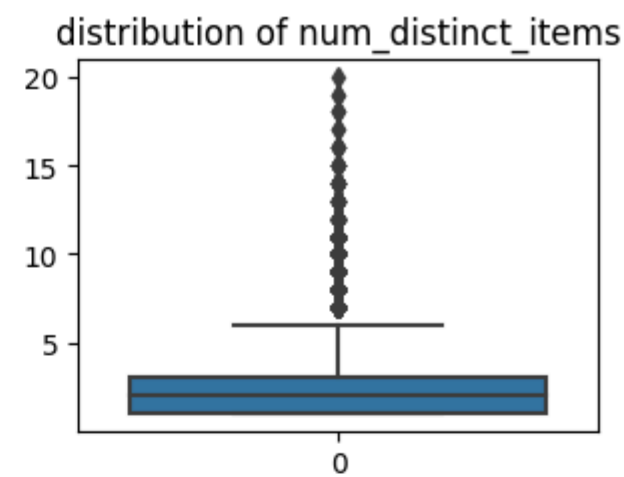
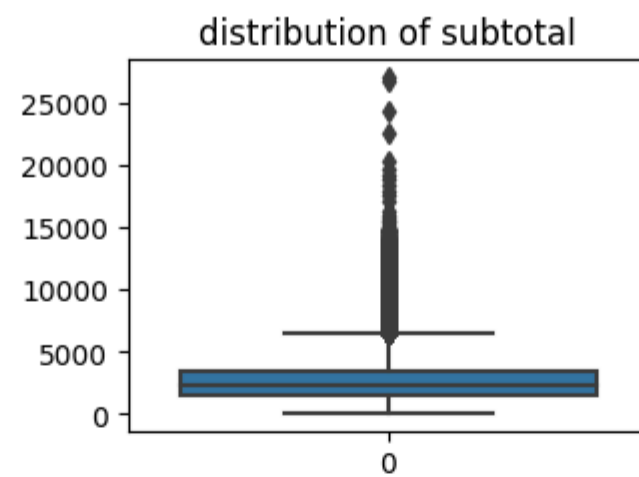
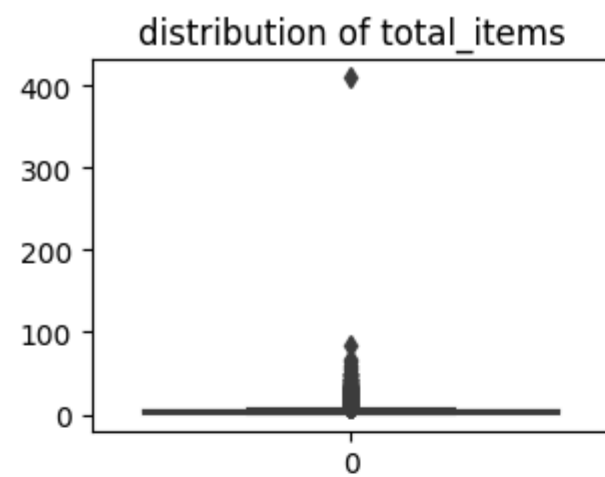
```
In [49]: 1 df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 195059 entries, 0 to 197427
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   market_id                             195059 non-null object
1   store_id                              195059 non-null object
2   store_primary_category                 195059 non-null object
3   order_protocol                         195059 non-null object
4   total_items                           195059 non-null int64
5   subtotal                              195059 non-null int64
6   num_distinct_items                    195059 non-null int64
7   min_item_price                        195059 non-null int64
8   max_item_price                        195059 non-null int64
9   total_onshift_partners                 195059 non-null float64
10  total_busy_partners                    195059 non-null float64
11  total_outstanding_orders                195059 non-null float64
12  time_taken_min                         195059 non-null float64
13  hour_created                           195059 non-null int32
14  day_of_week                            195059 non-null int32
dtypes: float64(4), int32(2), int64(5), object(4)
memory usage: 22.3+ MB
```

```
In [50]: 1 numerical_columns = df.select_dtypes(include=['int', 'float']).columns.tolist()
2 #numerical_columns= numerical_columns[2:]
3 numerical_columns
```

Out[50]: ['total_items',
'subtotal',
'num_distinct_items',
'min_item_price',
'max_item_price',
'total_onshift_partners',
'total_busy_partners',
'total_outstanding_orders',
'time_taken_min',
'hour_created',
'day_of_week']

```
In [51]: 1 fig = plt.figure(figsize=(10, 10))
2
3 for i in range(len(numerical_columns)):
4     plt.subplot(4, 3, i+1)
5     fig.tight_layout()
6     sns.boxplot( df[numerical_columns[i]])
7     plt.title(f"distribution of {numerical_columns[i]}")
```



```
In [52]: 1 df.loc[df["total_items"]>80]
```

Out[52]:

| | market_id | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_partners | total_outstanding_order |
|--|-----------|----------|----------------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|---------------------|-------------------------|
| | 47231 | 2.0 | c156cea027720c227089e679b3ae9d1b | fast | 411 | 3115 | 5 | 0 | 299 | 35.000000 | 35.000000 | 39.000 |
| | 182223 | 6.0 | e50372d3fee4eadec9c42aa6528097cc | fast | 84 | 1016 | 4 | 0 | 289 | 44.929771 | 41.896183 | 58.439 |

- As we can see there is only one record which have total_items more than 200

```
In [53]: 1 df= df.loc[df["total_items"]<80]
```

```
In [54]: 1 df.loc[df["max_item_price"]>10000]
```

Out[54]:

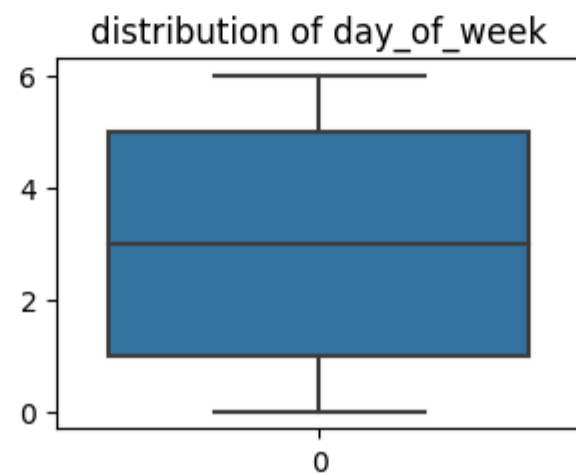
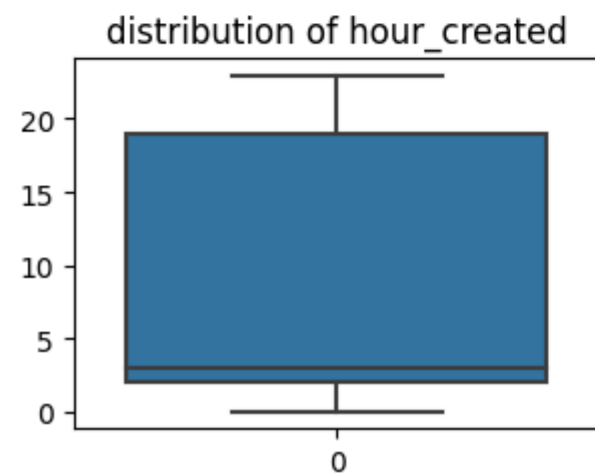
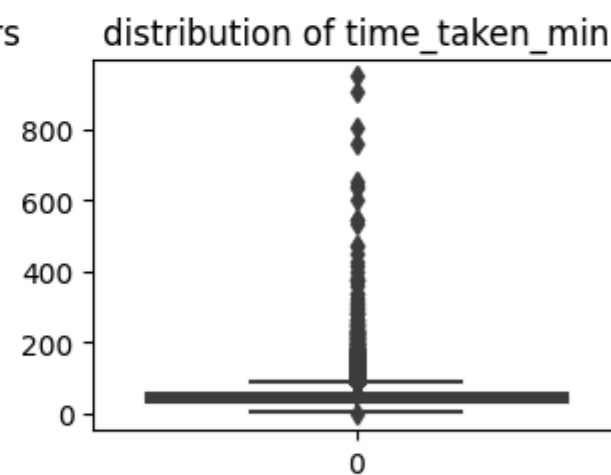
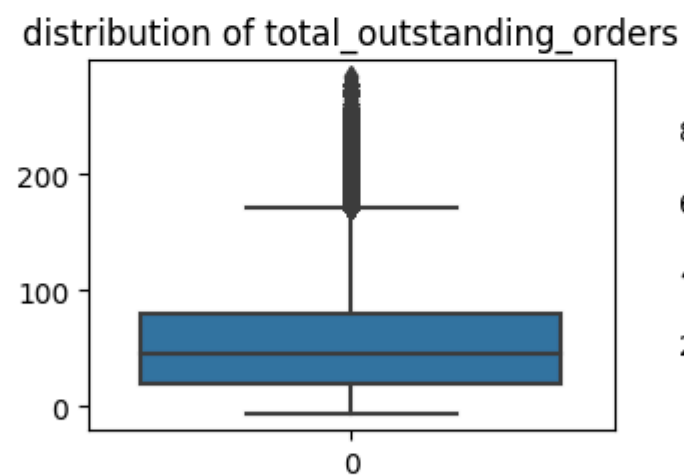
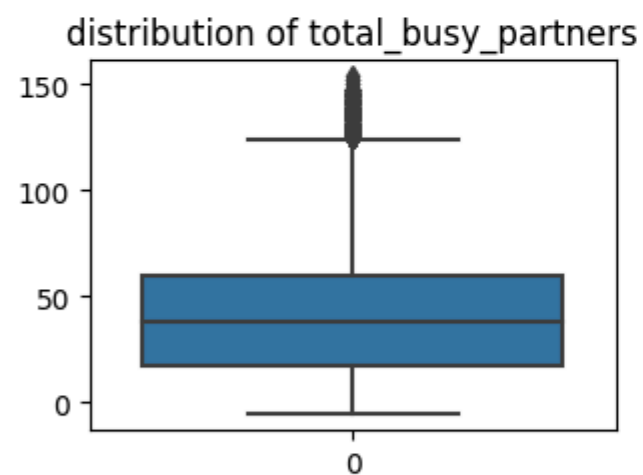
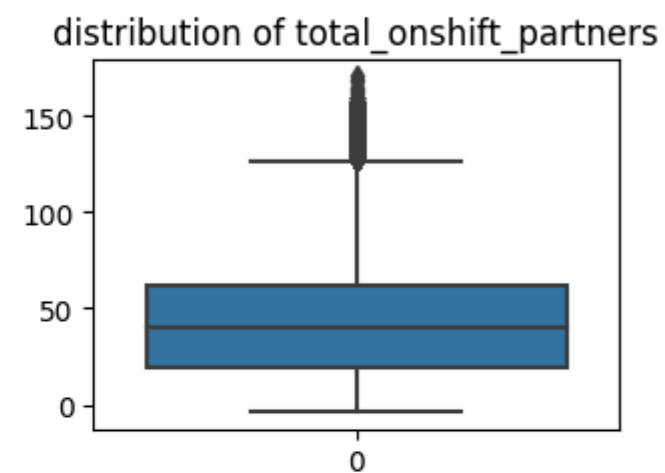
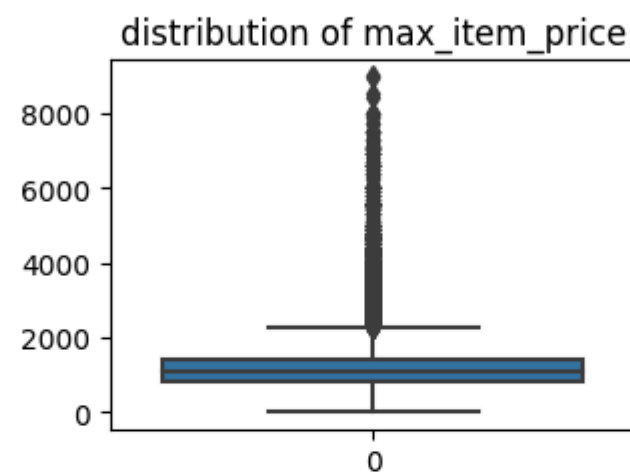
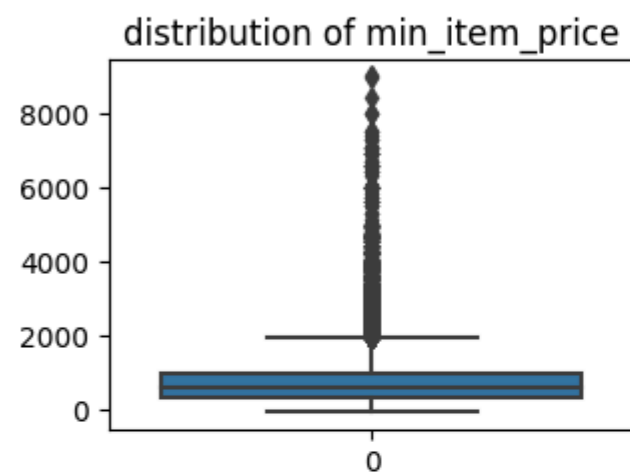
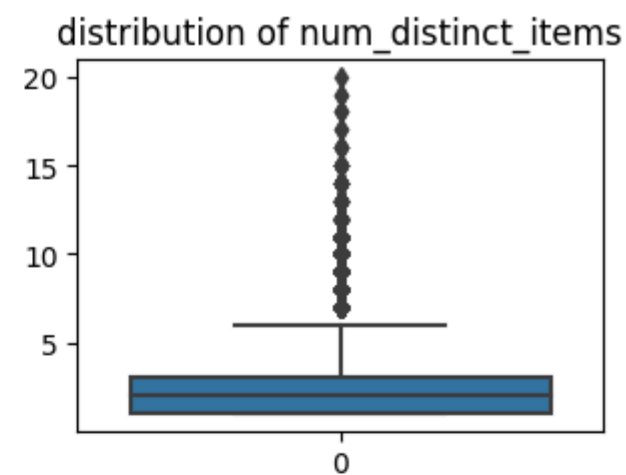
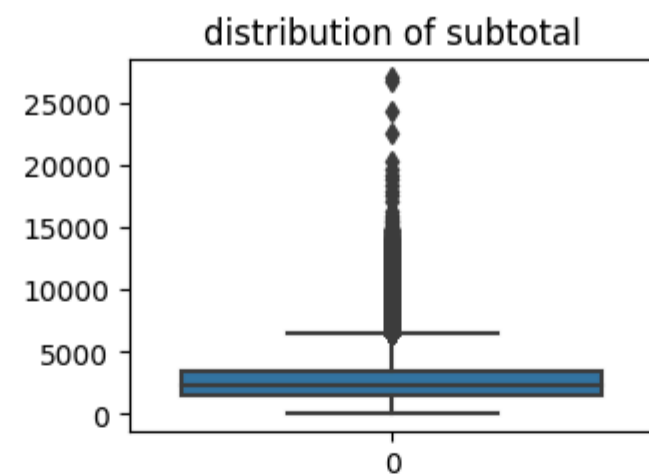
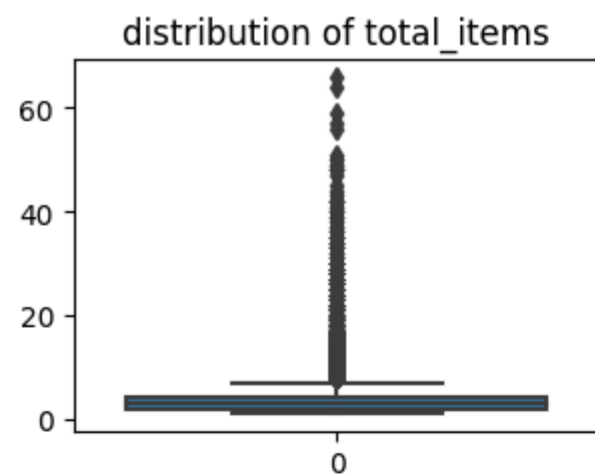
| | market_id | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_partners | total_outstanding_order |
|--|-----------|----------|----------------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|---------------------|-------------------------|
| | 139718 | 3.0 | 36452e720502e4da486d2f9f6b48a7bb | breakfast | 1 | 14700 | 1 | 14700 | 14700 | 23.0 | 21.0 | 2 |

```
In [55]: 1 df= df.loc[df["max_item_price"]<10000]
```

```
In [56]: 1 df.shape
```

Out[56]: (195056, 15)

```
In [57]: 1 fig = plt.figure(figsize=(10, 10))
2
3 for i in range(len(numerical_columns)):
4     plt.subplot(4, 3, i+1)
5     fig.tight_layout()
6     sns.boxplot( df[numerical_columns[i]])
7     plt.title(f"distribution of {numerical_columns[i]}")
```

```
In [58]: 1 df.loc[df["time_taken_min"]>600]
```

Out[58]:

| | market_id | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_partners | total_outstanding_order |
|--------|-----------|----------------------------------|------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|---------------------|-------------------------|
| 66787 | 6.0 | e85ca00d008a532279b798033d59a4c7 | italian | 1.0 | 1 | 795 | 1 | 795 | 795 | 44.929771 | 41.896183 | 58.439 |
| 76743 | 2.0 | 9380e398ee9bea45b992a3daaa6b7c4d | pizza | 6.0 | 1 | 990 | 1 | 795 | 795 | 114.000000 | 112.000000 | 184.000 |
| 83055 | 2.0 | c1502ae5a4d514baec129f72948c266e | burger | 4.0 | 3 | 2379 | 2 | 389 | 695 | 109.000000 | 102.000000 | 163.000 |
| 86952 | 3.0 | 831b342d8a83408e5960e9b0c5f31f0c | thai | 2.0 | 3 | 2185 | 3 | 495 | 995 | 19.000000 | 19.000000 | 16.000 |
| 105825 | 2.0 | 5fc7c9bd1fcb12799f02da8adfa4954f | alcohol | 5.0 | 3 | 2850 | 3 | 200 | 1500 | 96.000000 | 103.000000 | 156.000 |
| 139989 | 3.0 | a8e5a72192378802318bf51063153729 | mediterranean | 3.0 | 2 | 2300 | 2 | 600 | 1700 | 20.000000 | 19.000000 | 20.000 |
| 175971 | 1.0 | 5d616dd38211ebb5d6ec52986674b6e4 | mexican | 1.0 | 5 | 1530 | 3 | 300 | 315 | 24.000000 | 25.000000 | 30.000 |
| 190860 | 1.0 | b132ecc1609bfcf302615847c1caa69a | indian | 3.0 | 4 | 3660 | 4 | 375 | 1195 | 71.000000 | 70.000000 | 111.000 |

```
In [59]: 1 df= df.loc[df["time_taken_min"]<600]
```

```
In [60]: 1 # Independent variables
2 X= df.drop("time_taken_min", axis=1)
```

```
In [61]: 1 #target variable
2 y= df["time_taken_min"]
```

Splitting the data into Train, Validation and Test Data

```
In [62]: 1 X_train_val, X_test, y_train_val, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
2 X_train, X_val, y_train, y_val = train_test_split(X_train_val, y_train_val, test_size=0.2, random_state=42)
3
4 print('Train : ', X_train.shape, y_train.shape)
5 print('Validation:', X_val.shape, y_val.shape)
6 print('Test : ', X_test.shape, y_test.shape)
```

Train : (124830, 14) (124830,)
Validation: (31208, 14) (31208,)
Test : (39010, 14) (39010,)

Encoding categorical columns

```
In [63]: 1 enc = TargetEncoder(cols=["market_id", "store_id", "store_primary_category", "order_protocol"])
2 X_train = enc.fit_transform(X_train, y_train)
3
4 X_val = enc.transform(X_val)
5 X_test = enc.transform(X_test)
```

In [64]:

1 X_train.head()

Out[64]:

| | market_id | store_id | store_primary_category | order_protocol | total_items | subtotal | num_distinct_items | min_item_price | max_item_price | total_onshift_partners | total_busy_partners | total_outstanding_orders | hour_created | day_o |
|--------|-----------|-----------|------------------------|----------------|-------------|----------|--------------------|----------------|----------------|------------------------|---------------------|--------------------------|--------------|-------|
| 45046 | 46.099077 | 45.011699 | 47.334094 | 45.567002 | 1 | 1075 | 1 | 725 | 725 | 73.0 | 74.0 | 101.0 | 20 | |
| 87732 | 51.304882 | 49.318093 | 47.899950 | 49.918444 | 2 | 1150 | 2 | 295 | 855 | 20.0 | 15.0 | 16.0 | 3 | |
| 162823 | 46.429496 | 41.474057 | 46.279055 | 47.395853 | 5 | 1984 | 5 | 335 | 399 | 18.0 | 18.0 | 29.0 | 3 | |
| 159857 | 46.429496 | 44.946896 | 50.192517 | 48.594070 | 2 | 1143 | 2 | 279 | 729 | 14.0 | 12.0 | 12.0 | 23 | |
| 136409 | 46.099077 | 41.596319 | 45.531750 | 46.895346 | 4 | 1475 | 4 | 250 | 695 | 60.0 | 59.0 | 71.0 | 4 | |

Scaling the data for NN

In [65]:

1 scaler = StandardScaler()
2 X_train = scaler.fit_transform(X_train)
3
4 X_val = scaler.transform(X_val)
5 X_test = scaler.transform(X_test)

Keras Sequential API

Simple model

In [66]:

1 def baseline():
2 model= Sequential()
3 model.add(Dense(32,kernel_initializer='normal', activation= "relu", input_shape= (X_train.shape[1],)))
4 model.add(Dense(16, activation= "tanh"))
5 model.add(Dense(12, activation="relu"))
6 model.add(Dense(8, activation="relu"))
7 model.add(Dense(4))
8 model.add(Dense(1))
9 return model

In [67]:

1 simple_model= baseline()

In [68]:

1 simple_model.compile(optimizer="adam", loss="mean_squared_error", metrics=["mean_absolute_error"])

In [69]:

1

simple model summary

2

simple_model.summary()

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense) | (None, 32) | 480 |
| dense_1 (Dense) | (None, 16) | 528 |
| dense_2 (Dense) | (None, 12) | 204 |
| dense_3 (Dense) | (None, 8) | 104 |
| dense_4 (Dense) | (None, 4) | 36 |
| dense_5 (Dense) | (None, 1) | 5 |

Total params: 1,357 (5.30 KB)

Trainable params: 1,357 (5.30 KB)

Non-trainable params: 0 (0.00 B)

In [70]:

1

hist= simple_model.fit(X_train, y_train, epochs= 15, validation_data=(X_val, y_val), verbose=1)

Epoch 1/15
3901/3901 9s 2ms/step - loss: 469.1544 - mean_absolute_error: 15.0085 - val_loss: 264.8139 - val_mean_absolute_error: 11.2549

Epoch 2/15
3901/3901 8s 2ms/step - loss: 255.6663 - mean_absolute_error: 11.2893 - val_loss: 260.5962 - val_mean_absolute_error: 11.3811

Epoch 3/15
3901/3901 9s 2ms/step - loss: 251.8255 - mean_absolute_error: 11.1864 - val_loss: 258.9371 - val_mean_absolute_error: 11.2535

Epoch 4/15
3901/3901 9s 2ms/step - loss: 257.7316 - mean_absolute_error: 11.1421 - val_loss: 257.8633 - val_mean_absolute_error: 11.3856

Epoch 5/15
3901/3901 9s 2ms/step - loss: 254.7227 - mean_absolute_error: 11.1492 - val_loss: 257.4925 - val_mean_absolute_error: 11.5078

Epoch 6/15
3901/3901 9s 2ms/step - loss: 242.8393 - mean_absolute_error: 11.0477 - val_loss: 259.0905 - val_mean_absolute_error: 11.6076

Epoch 7/15
3901/3901 9s 2ms/step - loss: 241.5830 - mean_absolute_error: 11.0278 - val_loss: 255.7721 - val_mean_absolute_error: 11.2334

Epoch 8/15
3901/3901 8s 2ms/step - loss: 248.9611 - mean_absolute_error: 11.0652 - val_loss: 255.3586 - val_mean_absolute_error: 11.3797

Epoch 9/15
3901/3901 7s 2ms/step - loss: 245.3082 - mean_absolute_error: 11.0099 - val_loss: 256.0684 - val_mean_absolute_error: 11.3483

Epoch 10/15
3901/3901 7s 2ms/step - loss: 252.9208 - mean_absolute_error: 11.0220 - val_loss: 256.6858 - val_mean_absolute_error: 11.3409

Epoch 11/15
3901/3901 9s 2ms/step - loss: 245.4139 - mean_absolute_error: 11.0328 - val_loss: 256.6435 - val_mean_absolute_error: 11.1437

Epoch 12/15
3901/3901 9s 2ms/step - loss: 247.9652 - mean_absolute_error: 11.0007 - val_loss: 257.4635 - val_mean_absolute_error: 11.5273

Epoch 13/15
3901/3901 7s 2ms/step - loss: 245.4950 - mean_absolute_error: 10.9836 - val_loss: 258.7766 - val_mean_absolute_error: 11.2836

Epoch 14/15
3901/3901 7s 2ms/step - loss: 251.4392 - mean_absolute_error: 10.9753 - val_loss: 255.7395 - val_mean_absolute_error: 11.1996

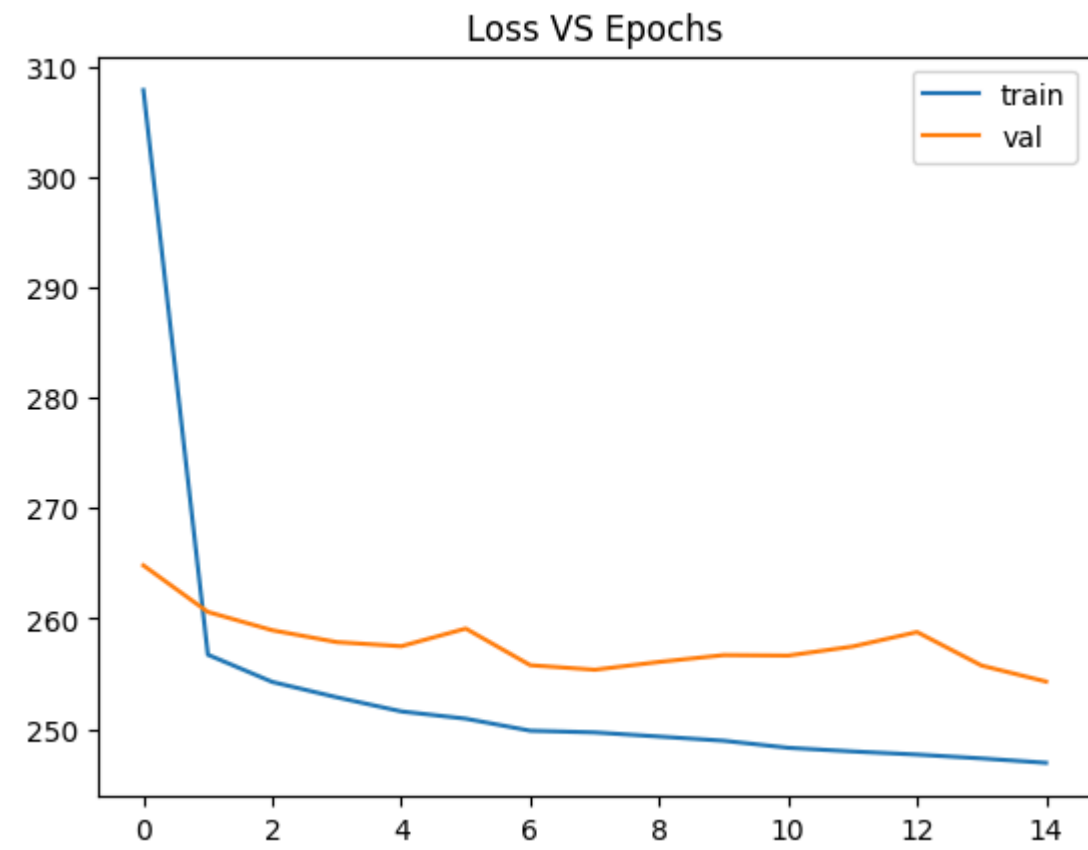
Epoch 15/15
3901/3901 7s 2ms/step - loss: 244.2113 - mean_absolute_error: 10.9364 - val_loss: 254.2893 - val_mean_absolute_error: 11.1425

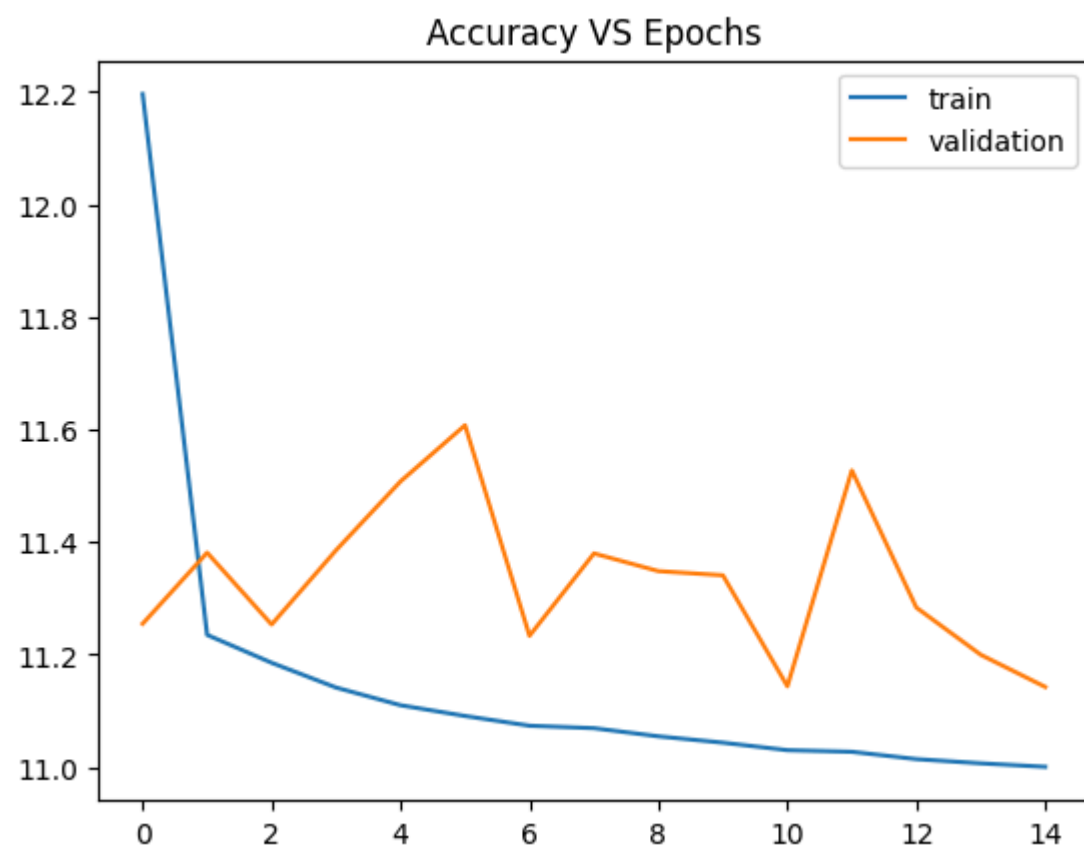
In [71]:

```
1 epochs = hist.epoch
2 train_MSE = hist.history["loss"]
3 train_MAE = hist.history["mean_absolute_error"]
4 val_MSE = hist.history["val_loss"]
5 val_MAE = hist.history["val_mean_absolute_error"]
```

In [72]:

```
1 plt.figure()
2 plt.plot(epochs, train_MSE, label="train")
3 plt.plot(epochs, val_MSE, label="val")
4 plt.legend()
5 plt.title("Loss VS Epochs")
6 plt.show()
7
8 plt.figure()
9 plt.plot(epochs, train_MAE, label="train")
10 plt.plot(epochs, val_MAE, label="validation")
11 plt.legend()
12 plt.title("Accuracy VS Epochs")
13 plt.show()
```





In [73]: 1 y_pred = simple_model.predict(X_test)

1220/1220 ————— 2s 2ms/step

In [74]: 1 r2_score_simple= r2_score(y_test, y_pred)

In [75]: 1 r2_score_simple

Out[75]: 0.26072215528845355

In []: 1

Using Keras Tuner

```
In [102]: 1 def build_model(hp):
2         model = Sequential()
3         model.add(Dense(units=hp.Int("units_layer1", min_value=128, max_value=1028, step=32),
4                               activation=hp.Choice("activation_layer1", ["relu", "tanh"]), input_shape= (X_train.shape[1],)))
5         model.add(BatchNormalization())
6
7
8         model.add(Dense(units=hp.Int("units_layer2", min_value=100, max_value=500, step=100),
9                               activation=hp.Choice("activation_layer2", ["relu", "tanh"])))
10
11        model.add(Dense(units=hp.Int("units_layer2", min_value=50, max_value=150, step=10),
12                              activation=hp.Choice("activation_layer2", ["relu", "tanh"])))
13
14        model.add(Dense(units=hp.Int("units_layer3", min_value=16, max_value=64, step=8),
15                              activation=hp.Choice("activation_layer3", ["relu", "tanh"])))
16        model.add(Dense(units=hp.Int("units_layer3", min_value=4, max_value=16, step= 1),
17                              activation=hp.Choice("activation_layer3", ["relu", "tanh"])))
18
19        model.add(Dense(1))
20
21        model.compile(
22            optimizer="adam",
23            loss="mean_squared_error",
24            metrics=["mean_absolute_error"])
25
26
27        return model
28
29
30 build_model(keras_tuner.HyperParameters())
```

Out[102]: <Sequential name=sequential_1, built=True>

```
In [103]: 1 tuner = keras_tuner.BayesianOptimization(
2         build_model,
3         objective='val_mean_absolute_error',
4         max_trials=10, overwrite=True)
```

```
In [104]: 1 tuner.search(X_train, y_train, epochs=10, validation_data= (X_val, y_val), callbacks=[EarlyStopping(patience=5, min_delta=80)], verbose= 1)
2 best_model = tuner.get_best_models()[0]
```

Trial 10 Complete [00h 02m 56s]
val_mean_absolute_error: 11.230664253234863

Best val_mean_absolute_error So Far: 11.16738224029541
Total elapsed time: 00h 14m 09s

```
In [105]: 1 best_model = tuner.get_best_models(num_models=1)[0]
2
```



```
In [106]: 1 best_model.fit(x=X_train, y= y_train, epochs=100, validation_data=(X_val, y_val), callbacks=[EarlyStopping(min_delta=80, patience=5)])
```

Epoch 1/100
3901/3901 ————— 9s 2ms/step - loss: 249.2085 - mean_absolute_error: 11.1513 - val_loss: 261.8080 - val_mean_absolute_error: 11.8153
Epoch 2/100
3901/3901 ————— 7s 2ms/step - loss: 245.4754 - mean_absolute_error: 11.1324 - val_loss: 261.8649 - val_mean_absolute_error: 11.3877
Epoch 3/100
3901/3901 ————— 7s 2ms/step - loss: 258.1833 - mean_absolute_error: 11.1644 - val_loss: 256.4775 - val_mean_absolute_error: 11.3045
Epoch 4/100
3901/3901 ————— 7s 2ms/step - loss: 245.5562 - mean_absolute_error: 11.0460 - val_loss: 254.7063 - val_mean_absolute_error: 11.1123
Epoch 5/100
3901/3901 ————— 7s 2ms/step - loss: 250.3231 - mean_absolute_error: 11.1135 - val_loss: 254.7024 - val_mean_absolute_error: 11.3584
Epoch 6/100
3901/3901 ————— 7s 2ms/step - loss: 248.9978 - mean_absolute_error: 11.0769 - val_loss: 255.6751 - val_mean_absolute_error: 11.0442

Out[106]: <keras.src.callbacks.history.History at 0x1c981d72950>

```
In [107]: 1 y_Pred_2= best_model.predict(X_test)
```

1220/1220 ————— 1s 985us/step

```
In [108]: 1 r2Score_2 = r2_score(y_test, y_Pred_2)
```

```
In [109]: 1 r2Score_2
```

Out[109]: 0.2563154199905652

```
In [110]: 1 best_model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense (Dense) | (None, 320) | 4,800 |
| batch_normalization (BatchNormalization) | (None, 320) | 1,280 |
| dense_1 (Dense) | (None, 100) | 32,100 |
| dense_2 (Dense) | (None, 100) | 10,100 |
| dense_3 (Dense) | (None, 32) | 3,232 |
| dense_4 (Dense) | (None, 32) | 1,056 |
| dense_5 (Dense) | (None, 1) | 33 |

Total params: 156,525 (611.43 KB)

Trainable params: 51,961 (202.97 KB)

Non-trainable params: 640 (2.50 KB)

Optimizer params: 103,924 (405.96 KB)

```
In [111]: 1 # MAPE Value
          2 from sklearn.metrics import mean_absolute_percentage_error
          3 mean_absolute_percentage_error(y_test, y_Pred_2)
```

```
Out[111]: 0.2410290383908038
```

- A MAPE of approximately 0.2618 means, on average, predictions have an error of approximately 26.18% relative to the actual values.

1. Defining the problem statements and where can this and modifications of this be used?

Problem Statement: The task is to build a neural network model to predict the delivery time of parcels using various features such as parcel weight, distance, delivery location, etc. The objective is to minimize the prediction error to ensure timely and efficient parcel delivery.

Applications and Modifications:

- **E-commerce:** Predicting delivery times for products to improve customer satisfaction.
- **Logistics:** Optimizing routes and schedules for delivery trucks.
- **Food Delivery Services:** Estimating delivery times for meals to ensure food arrives hot and fresh.
- **Courier Services:** Enhancing delivery speed predictions to manage customer expectations.
- **Supply Chain Management:** Predicting lead times to optimize inventory and reduce holding costs.

2. List 3 functions the pandas datetime provides with one line explanation.

- `pd.to_datetime()` : Converts argument to datetime.
- `pd.date_range()` : Generates a fixed frequency DatetimeIndex.
- `pd.DatetimeIndex()` : Provides an index of datetime64 data.

3. Short note on datetime, timedelta, time span (period)

- **datetime:** A module in Python for working with dates and times. It provides classes for manipulating dates and times in both simple and complex ways.
- **timedelta:** A class in the datetime module that represents the difference between two dates or times. It is used for calculating the difference and for adding or subtracting from a date or time.
- **time span (period):** Refers to a duration of time. In pandas, a Period represents a span of time defined by a frequency, such as a month, quarter, or year.

4. Why do we need to check for outliers in our data?

Outliers can significantly affect the performance of machine learning models. They can distort statistical measures, such as mean and standard deviation, and can lead to incorrect conclusions. By identifying and handling outliers, we can ensure that our model generalizes better and is not unduly influenced by anomalous data points.

5. Name 3 outlier removal methods?

- **IQR (Interquartile Range) Method:** Removes data points that fall below $Q1 - 1.5 \times IQR$ or above $Q3 + 1.5 \times IQR$.
- **Z-Score Method:** Removes data points that are a specified number of standard deviations away from the mean.
- **Isolation Forest:** An unsupervised learning algorithm that isolates observations by randomly selecting a feature and then randomly selecting a split value between the maximum and minimum values of the selected feature.

6. What classical machine learning methods can we use for this problem?

- **Linear Regression:** To predict delivery time based on input features.
- **Random Forest Regressor:** An ensemble method that uses multiple decision trees to improve predictive performance.
- **Support Vector Regression (SVR):** A regression technique that uses support vector machines to predict continuous values.

7. Why is scaling required for neural networks?

Scaling is required for neural networks to ensure that the input data is normalized, leading to faster convergence during training. It helps in making the gradient descent algorithm more efficient and stable. Scaling also prevents features with larger magnitudes from dominating the learning process.

8. Briefly explain your choice of optimizer.

Adam Optimizer: Adam (Adaptive Moment Estimation) is chosen because it combines the advantages of two other popular optimizers: AdaGrad and RMSProp. It maintains a per-parameter learning rate that improves performance on non-stationary problems and uses the first and second moments of the gradients to adapt the learning rate. This makes it suitable for a wide range of problems and is particularly effective for large datasets and high-dimensional parameter spaces.

9. Which activation function did you use and why?

ReLU and Tanh:

- **ReLU (Rectified Linear Unit):** Used because it helps in addressing the vanishing gradient problem, allowing the network to learn faster and perform better. It is computationally efficient and introduces non-linearity, which helps the network to learn complex patterns.
- **Tanh:** Used because it outputs values between -1 and 1, making it suitable for data that has been normalized. It also centers the data, which can lead to faster convergence.

10. Why does a neural network perform well on a large dataset?

In []:

| | |
|---|--|
| 1 | |
|---|--|